



Japan Information Platform Act YouTube Transparency Report

This report is published in accordance with Article 28 of the Act on the Limitation of Liability of Specified Telecommunications Service Providers for Damages and the Right to Demand Disclosure of Sender Identification Information (“Information Platform Act”).

Published: May 2026

Relevant Products/Service: YouTube

Company Name: Google LLC

Address: 251 Little Falls Drive, Wilmington, DE 19808, USA



Overview

This report is published in accordance with Article 28 of the Information Platform Act and provides data on our processes to review, remove and/or restrict access to content as well as statistical and categorical information on individual and government requests. The report also provides information about our content moderation practices and policies. This report covers the period from 26 July 2025 to 31 March 2026, except where indicated.

Introduction

At YouTube, our mission is to give everyone a voice and show them the world. Our policies, tools and processes help us maintain a responsible platform that viewers, creators and advertisers can rely on in order to continue to thrive on YouTube.

We remove content that violates our [Community Guidelines](#), and we have a process in place to comply with [local laws](#). Therefore, when content is removed from YouTube, it is either due to (1) legal reasons (“**Legal Removals**”), which generally apply to the relevant country or jurisdiction only; or (2) policy reasons (“**Policy Removals**”), which are global.

Here, we will first provide the number of monthly active users in Japan, followed by information on enforcement for Legal Removals and Policy Removals.

Monthly active users in Japan	106,400,000*
-------------------------------	--------------

* This represents the monthly average number of users who used the service while signed into their accounts.



Legal enforcement

Reporting illegal content

Users can report illegal content using [webforms](#) dedicated to specific legal issues such as trademark, copyright, counterfeit and defamation. Webforms may also be accessed via the in-product flagging feature after selecting “Legal issue” as the report reason. Our legal policies can be found on the Help Centre [here](#).

User applications for Legal Removals

In accordance with Article 28 (i) and (ii) of Information Platform Act, this data covers the volume of complaints filed by users requesting the removal or restriction of content.

Requests received for Legal Removals	Volume, Japan
Total	2,450
Circumvention	4
Counterfeit	110
Defamation	1,002
Other Legal	987
Privacy*	21
Trademark	326

* All removals under privacy guidelines are conducted through Legal Removals.

The categories shown above reflect what the user selected at the time of submission. We always send an immediate automated notification to requesters acknowledging receipt of their submission to the contact information provided.

Note that a single request may include multiple items with different reasons, or the reasons provided may differ from the legal grounds that ultimately led to the content's removal.

Notifications for Legal Removal requests (tickets) requiring extended review	Volume, Japan
Total	12
Defamation	4
Other Legal	8



An automated notification is sent for requests requiring extended review that remain pending seven days after the initial submission pursuant to Article 25 (2) of the Information Platform Act (the data point "Notifications for Legal Removal requests (tickets) requiring extended review" above represents this). If a single Legal Removal request (ticket) contains multiple items of content, this notification will be sent as long as at least one of those items remains unaddressed. Therefore, processing times are tracked and managed on a ticket basis, not per item of content.

As the requirements under Article 25 (2) (i) and (ii) of the Information Platform Act are not met, the reason provided in the notification for Legal Removal requests requiring extended review is classified as "unavoidable circumstances" under Article 25 (2)(iii). This includes cases where a request contains a large number of items or where the legal review is complex.

Although there are architectural constraints in our data aggregation system (e.g., the fact that multiple items of content may sit within 1 ticket), it is possible to monitor the status of expedited responses—which is the intent of Article 25—through the disclosure of the total number of “Requests received for Legal Removals” and the number of “Notifications for Legal Removal requests requiring extended review”.

Items identified for Legal Removal	Volume, Japan
Total	4,085
Not removed	3,796
Removed	289

Our systems are designed to send a prompt notification to the requestor when a definitive decision (Legal Removal or no action) is made on each requested item (videos, comments, or other associated entities) following a review by our specialised team. Note that the outcomes of several items may be contained within a single notification, or conversely, multiple notifications may occur for a single item.

Legal Removals by reason

Identifying violations of local law requires specific legal context and the submission of a formal assertion by a claimant. Consequently, all user requests for legal removal are initiated via reports for Legal Removals and are subject to manual review by professional review teams.

Note that while automated systems proactively detect potential violations of Community Guidelines, such systems are not used to identify or action content for legal violations.



YouTube evaluates Legal Removal requests based on legal grounds. While there are cases where our content policies overlap with local laws—e.g., CSAM—those removals are not considered legal removals and are recorded and reported through an entirely different infrastructure (our Community Guidelines flagging). The final reasons for Legal Removal and their respective numbers are as follows:

Legal Removal reason	Volume, Japan**
Total	289
Circumvention	0
Counterfeit	3
Defamation	155
Other Legal	13
Privacy*	1
Trademark	117

* All removals under privacy guidelines are conducted through Legal Removals.

** Items removed include videos, comments, and other entities (e.g. posts, thumbnails).

While users provide rationale for a request when they submit it (see the table of “Requests received for Legal Removals” above), we conduct our review without being strictly bound by this rationale. Our ultimate classification is determined and recorded on a per-item basis following internal expert legal review of the specific items cited. This is because a single submission may include multiple items with different rationales, rationales that are inconsistent with standard reporting categories, or rationales that are inaccurate and inconsistent with the legal reasons for which content is ultimately actioned (so a video may be removed for a different legal reason to that submitted by the user). Due to this structural difference, processing times are tracked on a per-ticket basis, whereas our decision rationales are recorded on a per-item basis. Consequently, our system architecture does not allow for the correlation of these distinct data sets.

Copyright enforcement

We offer a suite of copyright management tools that gives rights holders control of their copyrighted material on YouTube. These tools have allowed creativity to thrive, while giving creators and artists the tools and resources needed to manage their content online. The Copyright Transparency Report shares information about the copyright enforcement efforts on YouTube. More information on copyright enforcement is available on the Google Transparency Report website [here](#).



Government and court-ordered removal requests

This section is prepared in accordance with Article 28 (iv) of Information Platform Act. Governments contact Google, including YouTube, with content removal requests for a number of reasons. Government bodies may claim that content violates a local law, and include court orders that are often not directed at us with their requests. More information on government requests and court orders to remove content is available on the Google Transparency Report website [here](#).

We will continue to collaborate with MIC to provide content moderation data that meets the requirements of the Information Platform Act in future iterations of this report.

Policy enforcement

Community Guidelines

In addition to action taken on content in accordance with local laws, YouTube has a robust content moderation system consisting of content policies, called Community Guidelines; detection and enforcement processes; and a penalty and education system for content creators. The Community Guidelines policies are designed to address a range of harmful content, and enforcement systems aim to limit users' exposure to this policy-violating content. The Community Guidelines include, but are not limited to policies covering:

- [Spam, deceptive practices, and scams](#)
- [Impersonation](#)
- [Fake engagement](#)
- [Nudity and sexual content](#)
- [Violent or graphic content](#)
- [Harmful or dangerous content](#)
- [Suicide and self-harm policy](#)
- [Harassment and cyberbullying](#)

In addition, we aim to protect users' privacy and have established privacy guidelines regardless of local laws. You can find more information [here](#).

We continuously review and update policies to ensure that they reflect new types of harmful content. We share information about these [policy updates](#) through the Help Centre. Note that our [Terms of Service](#) cover the use of our service and incorporate our Community Guidelines.

Policy enforcement process

In general, enforcement of YouTube's policies is a joint effort between people and machine learning technology. We start by giving a team of experienced content moderators detailed



guidelines that explain what makes content violative and non-violative, and ask them to differentiate between violative and non-violative material. If the new guidelines allow them to achieve a very high level of accuracy, we expand the testing group to include moderators across different backgrounds, languages and experience levels.

Then we may begin revising the guidelines so that they can be accurately interpreted across a larger, more diverse set of moderators. These findings then help train our machine learning technology to detect potentially violative content at scale. As done with its content moderators, we also test our models to understand whether it has provided enough context for them to make accurate assessments about what to surface for people to review.

Once models are trained to identify potentially violative content, the role of content moderators remains essential throughout the enforcement process. Machine learning helps identify potentially violative content at scale and content moderators may then help assess whether the content should be removed. In some cases, there is high confidence that the content is violative given similar content that was previously removed.

This collaborative approach helps improve the accuracy of our models over time, as models continuously learn and adapt based on content moderator feedback. It also means our enforcement systems can manage the sheer scale of content that is uploaded to YouTube, while still digging into the nuances, where necessary, that determine whether a piece of content is violative.

Flagging

YouTube relies on a combination of people and automated systems to identify content that may violate the Community Guidelines and to enforce those policies. Content may first be flagged by automated systems, by members of the Priority Flagger programme, or by users in the broader YouTube community.

Automated flagging

YouTube's automated systems use machine learning that uses data from previous human reviews to identify potentially violative content. The systems are continuously supplied with millions of data points from human reviews, meaning these automated systems offer a high level of accuracy in detecting violations. Automated systems also provide efficient response time to users given the high volume of content that is uploaded to YouTube daily. The vast majority of content is first identified by automated systems, enabling us to enforce our policies quickly and at scale with a goal of limiting users' exposure to policy-violating content.

Human flags by flagger type

In addition to our automated flagging systems, YouTube offers users mechanisms to report or flag content that they believe violates our Community Guidelines or other legal issues.



Users can easily report content using our built-in flagging feature. Details on how to report different types of content using our flagging feature are outlined in YouTube's Help Centre. Users can visit their [Reporting History page](#) to check the status of videos they have reported on YouTube.

We also have a [Priority Flagger](#) programme, through which government agencies who are particularly effective at identifying certain types of harmful content have access to additional reporting processes and prioritised review of their flags. These flags are reviewed under the same policies and procedures as all other pieces of content we review.

Below are the total end-user flags from Japan IP addresses by reason.

Flagging Reason	User Flags, Japan
Total	7,693,292
Child Abuse	133,852
Harmful Dangerous Acts	683,070
Hateful or Abusive	1,470,212
Promotes Terrorism	236,531
Sexual	672,522
Spam or Misleading	2,861,555
Suicide, self-harm, or eating disorders	116,385
Violent or Repulsive	1,519,165

Review of flagged content

When systems have a high degree of confidence, they make an automated decision. If not, the system flags the content, and a trained human reviewer evaluates before any action is taken.

Employees that work on sensitive content moderation teams are offered subject matter specific training on a variety of topics. Employees working in sensitive content are required to complete a training on the Psychological Impact of Sensitive Content Review at the point of onboarding, and managers are required to complete an additional training on Supporting Teams who Work with Sensitive Content. Additional optional training opportunities include those on self-compassion, emotional agility,



and subject matter specific training to provide a deeper dive into the unique challenges faced by each team. The training is generally conducted via e-learning with opportunities for live facilitated training.

Policy Removals

In accordance with Article 28 (iv) of Information Platform Act, this data represents the total number of Policy Removals against videos, broken down by reason and by the source of the first detection.

Policy Removal reason	Volume of videos removed, Japan
Total	162,390
Child Safety	43,211
Harassment and Cyberbullying	70,695
Harmful or Dangerous	8,598
Hateful or Abusive	5,667
Misinformation	133
Nudity or Sexual	24,743
Other	215
Promotion of Violence and Violent Extremism	1,774
Spam, Deceptive Practices and Scams	1,007
Violent or Graphic	6,347

Source of first detection	Volume of videos removed, Japan
Automated detection	155,131
Government***	0
Organisation	1,652
User	5,607

*** Note the separate reporting above for Government Legal Removals



In our [Community Guidelines Enforcement Report](#), we provide data on 'Removed comments by the numbers' at a global scale. This information is updated quarterly. Due to our current data aggregation architecture, country-level granularity for action on comments is not feasible.

Notifications

This section is prepared in accordance with Article 28 (iii) of Information Platform Act.

Creator strikes

After a creator's first Community Guidelines violation, they will typically get a warning with no penalty to their channel. They will have the chance to take [policy training](#) to allow the warning to expire after 90 days. Creators will also get the chance to receive a warning in another policy category. If the same policy is violated within that 90 day window, the creator's channel will be given a strike.

If the creator receives three strikes in the same 90-day period, their channel may be removed from YouTube. In some cases, we may terminate a channel for a single case of severe abuse, as [explained in the Help Centre](#). We may also remove content for reasons other than Community Guidelines violations, such as a first-party privacy complaint or a court order. In these cases, creators will not be issued a strike.

Creator notifications

If a legal or policy enforcement action is taken against a video on a creator's channel or the channel itself, the creator will receive an email. A creator can also choose to receive notifications through in-app mobile and desktop notifications. The emails and notifications received by the creator explain the action taken on their content, information about the reason for action (whether local law or a specific YouTube policy), how it affects the user's channel, and what the creator can do next. More detailed guidelines of our processes and policies on strikes can be found [here](#).

We provide removal notifications for video actions and termination notices for channel-level actions, as these are our primary features, but we do not provide notifications for ancillary features, including comments. If we take comments as an example, we have various reasons for not sending the notification, including the fact that most removed comments can be removed for a number of reasons (including a creator moderating comments on their own video) and the majority of comments removed consist of high-volume deceptive commercial content (spam).

While our systems are designed to provide these notifications in all standard enforcement cases concerning the primary features, they may be withheld in exceptional circumstances where a justifiable reason exists. This primarily includes cases where notification is prohibited by legal authority or a court order, or where providing a notice poses a significant risk of secondary harm to a victim or complainant or would compromise the safety and integrity of enforcement systems against malicious actors. During this reporting period, there were zero instances in



Japan where notifications were withheld due to these justifiable reasons.

Restrictions (Service suspension measures)

In accordance with our Terms of Service, we reserve the right to restrict a creator's ability to create content on YouTube at our discretion. A channel may be restricted from using some or all YouTube features or may be fully terminated. Sometimes, a single case of severe abuse may result in [channel termination](#) without warning. If this happens, users are prohibited from using, creating, or acquiring another channel to get around these restrictions. This prohibition applies as long as the restriction remains active on the YouTube channel. A violation of this restriction is considered circumvention under [YouTube's Terms of Service](#), and may result in termination of all existing YouTube channels of the user, any new channels created or acquired, and channels in which the user is repeatedly or prominently featured.

This data reflects the number of cases where service suspension measures (i.e., channel terminations) were taken against Japanese users for violating Community Guidelines. We send notifications to all users subject to suspension.

Suspension Reason	Volume, Japan
Total	152,103
Child Safety	2,026
Harassment and Cyberbullying	375
Harmful or Dangerous	1,369
Hateful or Abusive	835
Misinformation	1,727
Multiple policy violations	110
Nudity or Sexual	6,000
Other	26
Promotion of Violence and Violent Extremism	177
Spam, Deceptive Practices, and Scams	139,405
Violent or Graphic	53



Appeals

User appeals process

The below appeals processes are available to users (creators) and are outlined in the YouTube Help Centre: [Appeal a Community Guidelines strike or video removal](#); [Appeal the age restriction of a video](#); [Appeal playlist or thumbnail removals](#); [Appeal a channel termination](#).

After a user submits an appeal, they will get an email from us letting them know the appeal outcome. One of the following will happen: If we find that a user's content followed YouTube's Community Guidelines, we will reinstate it and remove the strike from their channel. If a user appeals a warning and the appeal is granted, the next offence will be a warning. If we find that a user's content followed YouTube's Community Guidelines, but is not appropriate for all audiences, we will apply an age-restriction. If it is a video, it will not be visible to users who are signed out, are under 18 years of age, or have Restricted Mode turned on. If it is a custom thumbnail, it will be removed. If we find that a user's content was in violation of YouTube's Community Guidelines, the strike will stay and the video will remain down from the site. There is no additional penalty for appeals that are rejected.

Appeals against Policy Removals

This data reflects the number of videos appealed by Japanese users (creators) and the number of videos reinstated.

Action	Total videos, Japan
Appeals	20,759
Reinstatements	3,090

Japan's local context

This section is prepared in accordance with Article 28 (iv) of Information Platform Act.

YouTube recognises that a deep understanding of the local context specific to Japan, including cultural and social background, is essential to providing a safe and fair environment for Japanese users. We account for local context when it comes to events with specific impact for users in Japan, such as national elections or large-scale natural disasters.

Furthermore, in our standard content review process, we ensure that content moderators who understand the Japanese language are involved in reviewing Japanese-language content where it is vital to have accurate evaluation of linguistic nuances and context to review content.



The number of content moderators qualified to review videos in Japanese is as follows:

Qualified reviewers dedicated to content moderation - Japan	293
--	-----

In addition, for cases where additional legal expertise and professional judgement is specifically required due to the complexity of the matter or other factors, we have established a process to refer cases to external expert(s) (Infringing Information Investigation Specialists) in accordance with Article 24 (1) of the Information Platform Act. These expert(s) are qualified lawyers in Japan who are also well-versed in our policies and products. We make our final decisions after carefully considering their advice, which is based on their knowledge of Japanese law and judicial precedents. We trained the external expert(s) on our moderation processes to enable a complete and fair evaluation based on their expertise.

The number of cases where we sought the opinion of these experts during the period covered by this report is as follows:

Total expert investigation cases - Japan	6
---	---

Efforts supporting the Information Platform Act

This section is prepared in accordance with Article 28 (v) and (vi) of the Information Platform Act.

Over the years, we have worked continuously to develop policies, processes, and products to keep the YouTube community safe. As reflected in our Legal and Policy Removal processes, we are committed to keeping the service safe, for our users, advertisers, and society at large, while balancing open and free creative expression across the service.

Consistent with the spirit of the Information Platform Act, we have extensive controls that meet the requirements of the law in support of (1) user accessibility, (2) request status transparency, and (3) system effectiveness.

1. User accessibility

YouTube makes its content policies and Terms of Service available to users in clear, plain and intelligible language. We also operate a robust Help Centre that provides information about enforcement and reporting processes in easy-to-understand language. The method for submitting requests (Article 22 (1) of the Information Platform Act) can be easily accessed via YouTube's user interface and search services. Furthermore, pursuant to Article 21 of the Information Platform Act, our removal request contact points have been notified to the



Ministry of Internal Affairs and Communications (MIC) and are [listed here](#) along with those of other providers. We also established [this page](#) to make moderation policies and forms easy for users to find in one place.

In addition, Japan consistently ranks among the "Top 10 Countries/Regions with the Highest Volume of Flags (Reports)" data published in our [Community Guidelines Transparency Report](#). For example, in the quarter from October to December 2025, Japan ranked 7th globally in the volume of flags submitted by human users. We view this data as an indication that the reporting functionality provided by YouTube is known to users, readily accessed, and actively utilised by Japanese users. We also operate a Priority Flagger programme that provides government agencies with specific expertise in certain content policy areas a reporting pathway with prioritised review and additional communication channels, establishing an effective intake structure.

Finally, for creators who have been subject to enforcement actions, appeal entry points are accessible through the decision notification emails sent to creators, YouTube Creator Studio, and the YouTube Creator Studio Strikes summary.

2. Request status transparency

We work to strike a careful balance between providing clarity and transparency to users with the critical need to prevent malicious actors from circumventing and reverse engineering our criteria, processes, and systems. Under this principle, our standards for implementing measures to prevent transmission (Article 26 (1) of the Information Platform Act) are detailed in our Community Guidelines on [this page](#). In addition to the detailed information we provide in our Help Centre—which includes examples of the types of content that may violate our policies—we also have a dedicated website for creators with a convenient [Policies and Guidelines page](#) that makes it easy for users to find resources about our content policies and enforcement processes, including those related to legal issues like copyright.

Once a user makes a request, status transparency is built into our workflows. Upon receipt of a legal removal request, agents on the YouTube Legal Operations team process the request. For requests that are more complex and require calibration, they may escalate the request to an appropriate team member. Once a decision is finalised, a Legal Operations agent replies to the complainant with the final decision.

3. System effectiveness

YouTube's Violative View Rate (VVR) is our North Star metric that estimates the proportion of video views that violate its Community Guidelines (excluding spam). Voluntarily published quarterly since 2021, with data dating back to 2017, we closely monitor this metric to measure its progress on removing violative videos. In Q3 2025, VVR was 0.11-0.13% globally, and in Q4 2025, VVR was 0.14-0.15% globally. This means that out of every 10,000 views on YouTube in Q3, only 11-13 came from violative content and in Q4, only 14-15 came from



violative content. Our low VVR is achieved through our significant investment in systems that use a combination of machine learning detection and human moderators to detect infringing, violating, and problematic content at scale. As the absolute volume of content grows, we are committed to continuously updating and improving our detection systems.

We make available a qualitative description of VVR as the indicator of accuracy for automated means when used for content moderation. In 2021, we had our sampling methodology reviewed by MIT's Professor Arnold Barnett, who validated the approach [in this report](#). For more information, see our [Community Guidelines Transparency Report](#).

In addition, YouTube's Intelligence Desk monitors emerging trends related to potentially inappropriate content, and the YouTube Trust & Safety team partners with the Google Threat Intelligence Group (GTIG) to monitor and mitigate the impact of malicious actors around the globe at scale, including coordinated information operations that may affect Japan.

Finally, regarding request responsiveness, our support team ensures that inquiries received are monitored via a support dashboard to ensure that they are processed and prioritised for resolution within the standard business turnaround time. In addition, a sample of Cases are reviewed to ensure inquiries received are provided with a minimum "best efforts" support with responses.

We will continue to monitor external changes that influence our inherent risks and establish reasonable, proportionate, and effective measures to address risks. This will include measures that we can take alone and whole-of-society efforts that are best pursued collaboratively with other companies and stakeholders.