

Digital Services Act

Transparency Report for M365 Copilot

February 2025

Introduction

This report will cover M365 Consumer Copilot as it integrates with Microsoft 365 productivity apps (like Word, Excel, PowerPoint, Outlook, Loop, and more). Information in relation to Bing Image Creator, Bing Chat have not been addressed in this report

This report covers M365 Copilot consumer experiences (hereinafter “**Copilot**”). Copilot generally does not conduct “content moderation” as defined in the Digital Services Act, as content generated by Copilot, even when requested by users, is not content “provided by recipients of the service.” User queries and user prompts can trigger internal safety systems designed to ensure the services work as intended, including by preventing Copilot from generating content that could violate Microsoft policies or otherwise expose users to potentially harmful content. These activities do not qualify as “content moderation” as defined in Article 3(t) of the Digital Services Act since any outputs generated by Copilot (such as text, audio, or visual responses) are not content “provided by recipients of the service.” Nevertheless, we have provided overviews of how these systems operate within this report in the interests of transparency.

This DSA Transparency Report is responsive to the obligations under DSA Article 15(1), Article 24(1), and Article 42(2)-(3). This Report provides information regarding the following topics as they pertain to the European Union:

- Content Moderation at Microsoft’s Initiative
- Content Moderation & Automated Systems
- Content Moderation following a User Report
- Content Moderation Appeals
- Government Requests

Content Moderation at Microsoft's Initiative

Microsoft provides the information below in response to DSA Article 15(1)(c). This section reports data regarding content moderation for Copilot.

Use of AI-based classifiers

When using Copilot, a user may submit text, voice, images, and/or other enabled queries as inputs to the AI model, which then performs the relevant searches across the web (where applicable) and generates a response or, when asked to generate images, generates an image. This is known as a “prompt”. Copilot uses classifiers and content filters on user queries and prompts to help mitigate harm or prevent misuse. Examples include requests for information that could potentially lead to users being unexpectedly exposed to self-harm, violence, graphic content, hateful content, or misleading information. Flags from these classifiers may lead to mitigations, such as not returning generated content to the user or diverting the user to a different topic. Users remain free to seek information in other ways, such as using traditional web search, in the event an AI-based classifier restricts or limits a response.

User Terms Enforcement

Under the [M365 Consumer Copilot terms of use](#) and [Code of Conduct](#), violations may result in a permanent termination or permanent suspension of the user's access to Copilot.

Content Moderation & Automated Systems

Microsoft provides the information below in response to DSA Articles 15(1)(e) and 42(2)(c).

Content Safety:

Copilot relies on the amongst others [Azure OpenAI Service](#) for content filtering and moderation and leverages several internal tools including [Azure AI Content Safety](#). Azure AI Content Safety is an AI service that detects harmful user-generated and AI-generated content in applications and services. Azure AI Content Safety includes text and image APIs that allow detection of material that is harmful. The content filtering models for the Hate & Fairness, Sexual, Violence, and Self-harm categories have been specifically trained and tested in various languages. This system works by running both the input prompt and the response through classification models that are designed to identify and block the output of harmful content. By way of example, hate and fairness-related harms refer to any content that uses pejorative or discriminatory language based on attributes like race,

ethnicity, nationality, gender identity and expression, sexual orientation, religion, immigration status, ability status, personal appearance, and body size. Sexual content involves discussions about human reproductive organs, romantic relationships, acts portrayed in erotic or affectionate terms, pregnancy, physical sexual acts, including those portrayed as an assault or a forced act of sexual violence, prostitution, pornography, and abuse. Violence describes language related to physical actions that are intended to harm or kill, including actions, weapons, and related entities. Self-harm language refers to deliberate actions that are intended to injure or kill oneself.

Azure OpenAI Service's abuse monitoring system looks at usage patterns and employs algorithms and heuristics to detect and score indicators of potential abuse. Detected patterns consider, for example, the frequency and severity at which harmful content is detected (as indicated in content classifier signals) in a customer's prompts and completions, as well as the intentionality of the behavior. The trends and urgency of the detected pattern will also affect scoring of potential abuse severity. For example, a higher volume of harmful content classified as higher severity, or recurring conduct indicating intentionality (such as recurring jailbreak attempts) are both more likely to receive a high score indicating potential abuse. Prompts and completions that are flagged through content classification and/or identified as part of a potentially abusive pattern of use are subjected to another review process to help confirm the system's analysis and inform actioning decisions. By default, if prompts and completions are flagged through content classification as harmful and/or identified to be part of a potentially abusive pattern of use, they may be sampled for automated, eyes-off review by using an AI model and if necessary and required, human review. Except in cases of severe or recurring abuse, users typically are given an opportunity to explain or remediate and implement mechanisms to prevent recurrence of the abusive behavior.

For the period of this report, no EU content was actioned by Copilot. As noted above, user queries and user prompts can trigger internal safety systems designed to ensure the services work as intended, including by preventing Copilot from generating content that could violate Microsoft policies or otherwise expose users to potentially harmful content. These activities do not qualify as "content moderation" as defined in Article 3(t) of the Digital Services Act since any outputs generated by Copilot (such as text, audio, or visual responses) are not content "provided by recipients of the service."

For the period of the report, no policy violations were enforced.

Content Moderation following a User Report

Microsoft provides the information below in response to DSA Articles 15(1)(b)-(c) and 41(2)(a)-(b).

M365 Consumer Copilot does not have an in-product “report a concern” feature. This section is not applicable to M365 Consumer Copilot.

Content Moderation Appeals

Microsoft provides the information below in response to DSA Article 15(1)(d)

For the period of this report, Copilot did not receive any complaints that were appealed or reversed.

Government Requests

Microsoft provides the information below in response to DSA Article 15(1)(a)

For the period of this report, Copilot did not receive any requests from the EU member states.