

# Digital Services Act

## Transparency Report for Microsoft Advertising Powered by Xandr

February 2025

### Introduction

Microsoft Advertising Powered by Xandr, is a platform provider of digital and TV advertising solutions. It offers an end-to-end platform for executing campaigns on premium inventory across various screens. Xandr's technology focuses on data-enabled automation and audience solutions. The platform offers tools and services that enable advertisers, publishers, curators, and data providers to connect with target audiences, optimize advertising efforts, and achieve marketing goals.

This DSA Transparency Report is responsive to the obligations under DSA Article 15(1), Article 24(1), and Article 42(2)-(3). This Report provides information regarding the following topics as they pertain to the European Union:

- Content Moderation at Microsoft's Initiative
- Content Moderation & Automated Systems
- Content Moderation following a User Report
- Content Moderation Appeals
- Government Requests

### Content Moderation at Microsoft's Initiative

Microsoft provides the information below in response to DSA Article 15(1)(c). This section reports data regarding content moderation activities of the Microsoft Advertising platform powered by Xandr. These activities are taken at the Xandr platform's own initiative, absent a user report, to address content detected as violating Xandr Service Policies and Xandr Agreements.

Xandr content moderation includes the following mitigations:

- Xandr hosts and manages content provided by its advertisers, pursuant to customer agreements which include requirements to meet Xandr Service Policies. Xandr has clear and regularly enforced content policies and practices that prevent advertisements negatively impacting user safety. Microsoft Advertising powered by Xandr requires its advertisers and partners to comply with its policies throughout their use of the services. The Xandr Service Policies set out the requirements for ad content, including criteria upon which ad content will be prevented from serving. This includes prohibited content, restricted content policies, and intellectual property policies among others.
- The Xandr platform ensures that ads meet high quality standards by leveraging policies and procedures, which include both manual and automated processes to filter out and mitigate harmful online advertisement and traffic, including fraud and malware.
- Advertisements on the Xandr platform utilize a taxonomy that meets industry needs, automated category detection and a creative audit that includes automated and manual review processes to detect policy violations and risk. Detected prohibited content is banned and not made available to serve, detected restricted content is allowed based on specific restrictions based on that content and the remaining categorization enables blocking or permitting ad serving depending on the publisher preference. This offers publishers the ability to select delivery standards to allow/disallow categories of ads as eligible to serve on their properties. In addition, we incorporate these signals into our auction decisioning to support market level restrictions.
- We use a combination of automated and manual review categorization checks for risk and policy compliance with the Xandr Service Policies. This review leverages automated screening, machine-learning techniques, and the expertise of its ad quality operations team.
- Advertisers retain ownership of and responsibility for their ad content.

The Xandr platform can restrict or prevent delivery of ads, domains, accounts and customers based on the prohibited content, restricted content and intellectual property policies among others.

The Xandr platform uses automated content classifiers and manual detection on ads components and ad creatives when ads are created.

During the creative audit process, whether by human auditor or automation, ad creatives are assigned a brand and associated brand category, language, technical attributes, as well

as relevant creative-level sensitive category flags. Each of these labels may be allowed or disallowed by publishers in accordance with their publisher specific preferences.

If an ad creative is found to be in violation of the Xandr Service Policies, the ad creative is banned from serving across the platform.

We have models for detecting banned content as well as the sensitive categories in publisher ad quality profiles. Additionally, we have automated models for detecting brands, some of which are associated with sensitive categories as well.

The models will reject creatives when they have high confidence that an ad creative contains banned content and will send the creative for a secondary human review when banned content is detected with lower confidence. A similar strategy is used for sensitive categories, with the category being applied directly when models have high confidence and sent to human review to the extent possible when a model has lower confidence.

Scanner automation is also applied to detect and act on the following behaviors: security compliance, malware in the ad creative, presence of a malicious or fraudulent domain, technical behavior violations such as triggering an unexpected call chain of web links and previously identified high risk domains detected in ad creative or landing page. The scanner tools perform checks on technical attributes, malware and domain. This service also extracts assets including a creative screenshot, creative text, landing page screenshot and text which assists in decisioning and are passed on the automated classifiers and/or manual creative audit for action.

#### Product Data: Microsoft Advertising Powered by Xandr

| Policy Violation                                 | Number of measures taken at the provider's own initiative | Number of measures taken after detection with solely automated means |
|--------------------------------------------------|-----------------------------------------------------------|----------------------------------------------------------------------|
| IP Infringements                                 | 1                                                         | 1                                                                    |
| Negative Effects on Civic Discourse or Elections | 12,729                                                    | 12,729                                                               |
| Pornography or Sexualized Content                | 3,188                                                     | 3,188                                                                |
| Scams and Fraud                                  | 1,394                                                     | 1,394                                                                |
| Unsafe and Prohibited Products                   | 43,090                                                    | 43,090                                                               |

## Content Moderation & Automated Systems

Microsoft provides the information below in response to DSA Articles 15(1)(e) and 42(2)(c).

The Xandr platform has models for detecting banned content as well as the sensitive categories in publisher ad quality profiles. Additionally, we have automated models for detecting brands, some of which are associated with sensitive categories as well. If the landing page is rejected, creative automatically fails the creative audit.

The models will reject creatives when they have high confidence that an ad creative contains banned content and will send the creative for a secondary human review when banned content is detected with lower confidence. A similar strategy is used for sensitive categories, with the category being applied directly when models have high confidence and sent to human review to the extent possible when a model has lower confidence.

Scanner automation is also applied to detect and act on the following behaviors: security compliance, malware in the ad creative, presence of a malicious or fraudulent domain, technical behavior violations such as triggering an unexpected call chain of web links and previously identified high risk domains detected in ad creative or landing page. The scanner tools perform checks on technical attributes, malware and domain. This service also extracts assets including a creative screenshot, creative text, landing page screenshot and text which assists in decisioning and are passed on the automated classifiers and/or manual creative audit for action.

In addition to the automated content moderation processes, the Xandr platform moderation services are informed by utilizing signals from the wider Microsoft risk mitigation ecosystem to inform malvertising detection.

Content moderation and risk detection models are built using human labeled data to train the models and in the validation process to measure model prediction quality. The models are a combination of pretrained LLMs as well as fine-tuned machine learning neural networks. If the automated model is unable to categorize with high confidence, the ad creative is submitted for manual creative audit.

We have a dedicated number of human reviewers that label the classification models in sensitive category areas and prohibited content. There is also a large set of dedicated human review resources that supports the manual creative audit.

As category model training occurs, we separate out a portion of the human labeled data to act as our comparison baseline. This comparison helps compute performance metrics based on the human label data set to make the determination for that particular category if our system will use an automation action or seek human review for manual action. This precision scoring determines both when we release automation and informs when additional tuning is needed to drive continuous improvement into the model.

We also use frequency modeling using signals such as landing page domain association with frequently banned creatives to inform future automated banning action associated with that landing page.

Automation in content moderation consists of the following categories using detection models and human review solutions to ban prohibited content and prevent ads from serving (list not exhaustive):

Banned for prohibited content:

- Illegal content
- Hate speech
- Violent, morally reprehensible or defamatory content
- Illegal drugs, pharmaceuticals, or drug paraphernalia
- Adult content
- Malware
- Copyright and intellectual property violations
- Weapons content
- Extreme religious content
- CBD content
- Cryptocurrency products and services
- Political advertising
- Tobacco and vape content
- Ad behavioral violations such as executing files without intentional user interaction or displaying fake errors to induce a user action
- Information integrity, deceptive and misleading content violations
- Third party government services

The accuracy of the automation is tested offline by comparing the performance on a manual (human audit) labeled test dataset. In addition, there are checks performed on a subset of ads that are served on the platform.

As shared above, when category model training occurs, we separate out a portion of the human labeled data to act as our comparison baseline. This comparison helps compute performance metrics based on the human label data set to make the determination for that particular category if our system will use an automation action or seek human review for manual action. This precision scoring determines both when we release automation and informs when additional tuning is needed to drive continuous improvement into the model.

In addition, customers may escalate for remedy through an appeals process which is a manual and human review supported process.

The measures solely taken by automated means value includes edge case measures where some actions were taken but not registered as the right level as automated. We estimate that such number is less than 100 events.

|                                                    | Microsoft Advertising Powered by Xandr |
|----------------------------------------------------|----------------------------------------|
| Number of measures not taken by automated means    | 40,921                                 |
| Number of measures solely taken by automated means | 19,481                                 |

## Content Moderation following a User Report

Microsoft provides the information below in response to DSA Articles 15(1)(b)-(c) and 41(2)(a)-(b).

Microsoft Advertising supports Xandr by providing access for users to report ads they believe violate compliance associated with a particular harm through using the Microsoft Advertising Report a Concern form.

- When users have specific concerns about ads they see, they may report concerns through the Microsoft Advertising Report a Concern tool. This method escalates concerns to a team of reviewers for compliance assessment. Users can report:
  - Malvertising
  - Disallowed Content
  - Relevancy Concerns
- Users can report relevance, content or other issues including:
  - Disallowed Content: Bullying or harassment, Graphic violence and human gore, Threat or exposure of private sexual images, Child sexual abuse, Hate speech, Suicide or self-harm, offering or requesting sex, terrorism or violent extremism, coordinating to physically harm, threats or praise of violence, trafficking, fraud, phishing or scams, viruses, spyware or malware, infringement and other concerns not listed.
  - Relevance: ad is not relevant, landing page is not relevant, ad copy does not make sense, display URL does not match the landing page or other relevance issues.

- Page or site quality: high percentage of ads or links on the landing page, low value, sparse or limited content across the site, site redirection to an unrelated location or domain, other page or site quality issues or spam.
- Personally identifiable information (PII): site asks for unexpected personal information or phishing.
- Landing page navigation: site changes browser preferences without consent, site spawns multiple pop ups or prevents leaving the site, landing page does not load, product not available messaging or other landing page navigation issues.
- Users may also report ads they believe may infringe on intellectual property rights through the Intellectual Property Complaint form.
- Microsoft Advertising conducts a manual review of advertisements flagged to its customer support team and removes advertisements that violate its policies. Users will receive notification after action is taken.

## Content Moderation Appeals

Microsoft provides the information below in response to DSA Article 15(1)(d)

The appeals process on the Xandr platform does not currently use a tool that can track appeals based on the customer and account management engagement model. We support appeals on decisions and, if appropriate, act to overturn such decisions when made aware of by the account team. However, these actions are not logged in the tool in such a way that they are available to report. We are taking steps to update the process by enabling a ticketing solution that customers and account managers will use to initiate appeals, track actions and support reporting in the future.

## Government Requests

Microsoft provides the information below in response to DSA Article 15(1)(a)

Microsoft Advertising powered by Xandr did not receive any Government Requests in this time period.