

# Digital Services Act

## Transparency Report for Microsoft Advertising

February 2025

### Introduction

Microsoft Advertising is a comprehensive digital advertising platform designed to help businesses reach their target audiences and achieve their advertising goals. The platform allows businesses to drive sales by targeting users actively searching for their products or services. These ads appear on Bing search engine results pages and other web properties, helping businesses connect with potential customers at the moment they are looking for relevant information.

This DSA Transparency Report is responsive to the obligations under DSA Article 15(1), Article 24(1), and Article 42(2)-(3). This Report provides information regarding the following topics as they pertain to the European Union:

- Content Moderation at Microsoft's Initiative
- Content Moderation & Automated Systems
- Content Moderation following a User Report
- Content Moderation Appeals

### Content Moderation at Microsoft's Initiative

Microsoft provides the information below in response to DSA Article 15(1)(c). This section reports data regarding content moderation Microsoft Advertising undertakes at its own initiative, absent a user report, to address illegal content or information in violation of Microsoft Advertising's terms and conditions, which can be found in the Microsoft Advertising Agreement.

Microsoft Advertising service is provided under the Microsoft Advertising Agreement. The Microsoft Advertising Agreement was last updated in February 2024. One of the updates is to Section 2—Content Usage and Moderation, which states that Microsoft may need to (i) reproduce and publicly display, and (ii) disclose, as required by law or government and/or regulator orders, certain information relating to the counterparty's use of the services.

Microsoft further clarified that Microsoft reserves the right to review and restrict the counterparty's content per the Microsoft Advertising Policies. A full list of updates to the Microsoft Advertising Agreement can be viewed at the Microsoft Advertising agreement change log.

Microsoft Advertising content moderation includes the following mitigations:

- Microsoft Advertising hosts and manages content provided by its advertisers, pursuant to the Microsoft Advertising Agreement, which includes the Microsoft Advertising Network Policies. Microsoft Advertising has clearly and regularly enforced content policies and practices that prevent advertisements negatively impacting user safety. The Microsoft Advertising Policies set out the requirements for ad content, including criteria upon which ad content will be rejected, preventing ad serving. Microsoft requires its advertisers and partners to comply with its policies throughout their use of the Microsoft services. The Microsoft Advertising Network Policies set out the requirements for ad content, including criteria upon which ad content will be removed. This includes Disallowed Content policies, Restricted Content policies, Intellectual Property policies among others. Microsoft Advertising also has a set of Relevance and Quality Policies to manage the relevance and quality of the advertisements that it serves through its advertising network.
- Microsoft Advertising ensures that ads meet high relevance and quality standards by leveraging the Microsoft Advertising policies and procedures, which include both manual and automated processes to filter out and mitigate harmful online advertisement and traffic, including fraud and malware.
- Advertisements on the Microsoft Advertising platform are reviewed for compliance with the Microsoft Advertising Network policy prior to being posted. This review leverages automated screening including machine-learning, the expertise of its operations team, and dedicated user safety experts.
- Advertisers retain ownership of and responsibility for their ad content, but the advertiser must agree to the terms when signing up for a Microsoft Advertising account.
- The Microsoft Advertising platform can restrict or block ads, domains, accounts and customers based on the Microsoft Advertising Network Disallowed Content policies, Restricted Content policies, Intellectual Property policies among others.

Microsoft Advertising uses automated content detection on keywords, ads components, ad creatives and customer information when new customers sign up, new accounts and campaigns are created, new ads are created and when there are changes to these areas.

Microsoft Advertising uses machine learning models and techniques, large language model detection, automated screening, rule sets, the expertise of its operations team and user safety experts. This includes the following methods:

- Microsoft Advertising employs dedicated operational support and engineering resources to enforce these policies, combining various enforcement methods to prevent or take down advertisements that violate its policies.
- Microsoft uses both manual and systemic methods to review ads for compliance with their policies and reserves the right to remove those that are non-compliant.
- Microsoft Advertising continually seeks to add new signals and information sources to improve its detection systems, evaluating the incremental performance of adding these signals, and sources to sweeps and models and adding those that it finds useful to the system.
- Understanding the critical importance of maintaining a trustworthy ad ecosystem, Microsoft Advertising enforces content policies that strictly prohibit misleading, deceptive, and fraudulent advertisement content. This is complemented by a set of Relevance and Quality Policies aimed at ensuring the served advertisements are not only legal but also relevant and of high quality, thus preventing advertisers from utilizing questionable or misleading tactics to lure users. Advertising employs a combination of human and automated review process using algorithm systems and a robust filtration system to detect and mitigate harmful online traffic such as fraud, phishing, or malware and users can report advertising for takedown through the Bing “feedback” solution as well as Report a Concern form entry points. Detection on violative ad content covers ad text, images, extensions, product feeds and components, landing pages, URLs, keywords, and targeting settings. Malware specifically is considered an egregious violation and will result in the suspension of an advertising account.
- The Microsoft Advertising consumer protection and fraud risk detection system is multi-tiered, with checks based on main line machine-learning models, heuristic sweeps/rules, and manual reviews to prevent bad actors from coming onto the system. Overall, Microsoft Advertising has models and rules that apply at the advertiser-level as well as those that apply to a particular ad.
  - Some of the models that apply at the advertiser-level include:
    - Payment fraud detection: This model triggers when the advertiser adds a payment instrument. This model is owned by the commerce risk team and Microsoft Advertising consumes its decision.

- Purchase risk model: This model triggers upon billing the advertiser to determine whether to try to make the charge or block the advertiser. This helps to bring up-to-date information to determine payment risk.
  - Behavioral risk model: This model looks at various features of an advertiser's account or campaign(s) such as the ad text, keywords, account age, attributes of the payment instrument, individuals associated with the account, etc. to determine the risk of an advertiser being fraudulent. This model runs when the advertiser makes any change to their ad campaigns.
  - LLM-based spoofing detection and phishing prevention: This LLM-based approach works to predict if an advertiser is trying to deceive users into believing that an ad leads to a legitimate site or official website.
- Based on the above models' output, Microsoft Advertising decides to either suspend the advertiser or send it to manual review for a decision. In addition to the models described above, the platform evaluates and uses other signals.
- The Microsoft Advertising content moderation detection system is also multi-tiered, with checks based on classifier models and manual rule sets. Based on the model's output, Microsoft Advertising decides to reject or send content for manual review for a decision. Manual reviews are supported to detect policy compliance of ads when automated rejection does not have sufficient confidence. Content moderation also has models and rules that apply at the advertiser-level as well as those that apply to a particular ad.
  - The Microsoft Advertising platform applies models to individual ads based on the corresponding Microsoft Advertising Policies applicable to the market where the ads serve. Such models include, among others: adult, sensitive topics (hate speech, violence, gore), gambling.
  - There are classifiers applied based on certain policy categories, extensive list specific topics that ad copy is filtered against and review processes to ensure low automation confidence decisions are correct.
  - Actions are taken at the keyword level, ad component, ad level, and account or customer level.

## Product Data: Microsoft Advertising

| Policy Violation                                 | Number of measures taken at the provider's own initiative | Number of measures taken after detection with solely automated means | Account restriction Suspension |
|--------------------------------------------------|-----------------------------------------------------------|----------------------------------------------------------------------|--------------------------------|
| IP Infringements                                 | 497,239,847                                               | 497,239,847                                                          |                                |
| Negative Effects on Civic Discourse or Elections | 33,525                                                    | 33,525                                                               |                                |
| Pornography or Sexualized Content                | 1,233,525,478                                             | 1,233,525,478                                                        |                                |
| Scams and Fraud                                  | 66,561,647                                                | 66,561,647                                                           | 7,311                          |
| Scope of Platform Services                       | 353,387,867                                               | 353,387,867                                                          | 21,331                         |
| Unsafe and Prohibited Products                   | 14,572,988,333                                            | 14,572,988,333                                                       |                                |

## Content Moderation & Automated Systems

Microsoft provides the information below in response to DSA Articles 15(1)(e) and 42(2)(c).

Microsoft Advertising uses machine learning models and techniques, large language model detection, automated screening, rule sets, the expertise of its operations team and user safety experts.

- Microsoft applies robust classifiers to scan ads and websites promoted by advertisers for content policy compliance, misleading behavior, malware, phishing, fraudulent links, and other non-compliant behavior.
- Microsoft Advertising employs a robust filtration system to detect robotic traffic and other harmful cyber-attacks.
- Egregious violations of our terms or policies may warrant the immediate suspension of an advertiser.
- Microsoft Advertising has several teams of security engineers, support agents, and traffic quality professionals dedicated to continually developing and improving this traffic filtration and network monitoring system.
- Microsoft Advertising's support teams work closely with its advertisers to review complaints around suspicious online activity, and across internal teams to verify data accuracy and integrity.
- Any ads and associated components within an advertiser's account which violates policies within the following categories (Ad Requirements, Disallowed Content, Extensions, IP, Legal and Privacy, Media Formats, Product Ads, Relevance and Quality, Restricted Content), will be subject to a "three strikes" enforcement penalty

before Microsoft Advertising will suspend the user's account use of Microsoft Advertising services. If an advertiser is under a strike penalty, they will not be able to log into their account through Editor or the Mobile app but can access through the Web User Interface (UI). Microsoft Advertising determines strikes by policy categories related to the violation, or any violation that may pose a risk to the safety or security of Bing customers or users. Enforcement actions may include suspension of ads and associated components of the advertiser, limiting creation of new accounts, and suspending the advertiser's account temporarily or permanently.

- Microsoft Advertising and specialized cross-Microsoft teams regularly engage with external stakeholders in areas of key policy priorities, such as terrorist and violent extremist content, information integrity and misinformation, child sexual exploitation and abuse imagery (CSEAI), responsible AI, and AI-specific risks, hate speech, minors and technology, and copyright and trademark infringement, to ensure that its internal policies, practices, and standards are addressing key concerns. These engagements can inform processes of identifying, assessing, and mitigating risks across Microsoft Advertising and provide greater understanding of concerns related to the ads and broader societal and technology related issues.

Microsoft Advertising risk and content moderation models use human labeling to support precision, recall and confidence rates. This includes human labeling for initial training and continuous maintenance through audits to measure performance. Accuracy tuning is enabled by ongoing human labeling training efforts and through using human review decisions in rejections and overturned appeals to support ongoing adjustment to real-time issues.

Automation in content moderation consists of the following categories using content classifiers, term list management for automated filtering of ads or both (list not exhaustive):

- Platform and User Safety – abuse of platform and services including abuse of our users:
  - payment instrument fraud
  - circumvention or evading platform safeguards
  - misleading, clickbait and deceptive activity
  - fraudulent links, spam and scams
  - malware, scareware, ransomware
  - phishing
  - cloaking and spoofing detection

- repeated violations
- Other Disallowed Policy – disallowed content, products or services:
  - illegal content, products and services – including content flagged by wider Microsoft services such as CSEAI, nonconsensual intimate imagery and intimate extortion, terrorism or violent extremism. If Microsoft Advertising becomes aware of content containing apparent violations of these topics, we report to our Digital Safety team which ensures reporting to the proper authorities.
  - areas of questionable legality
  - deceptive products and services
  - endangered or threatened species
  - fireworks or dangerous materials
  - information integrity and misleading content
  - disallowed software behaviors
  - offensive, harmful or violent content including bullying and harassment, coordination of harm, graphic violence and human gore, hate speech, suicide or self-injury, violent threats, incitement and glorification of violence.
  - piracy
  - political advertising
  - sensitive advertising (tragic events)
  - tobacco products
  - trafficking or exploitation including sexual solicitation
- Intellectual property
- Restricted Policy - content that may be restricted in some manner or disallowed in a particular market:
  - adult content
  - alcohol
  - beauty and cosmetics
  - dating and personals
  - finance
  - food products
  - gambling and lotteries
  - government services and products
  - legal services
  - pharmacy and healthcare
  - religious content
  - ticket reselling
  - weapons

- Ad copy and landing page behaviors

Microsoft Advertising uses a combination of manual and automated processes to identify quality of actions, allow for remedy through a scaled appeals process, respond rapidly to new scenarios through list filter creation, monitor automated decisions by trending appeals rates for incorrect rejections and allow system and list decision reversal for manually reviewed mitigations to correct results.

In addition to the above, we utilize the “three strikes” enforcement penalty to remove repeat offenders and to protect the network at large. This process is described in section 3 above.

|                                                    | Microsoft Advertising |
|----------------------------------------------------|-----------------------|
| Number of measures not taken by automated means    | 2,032,159             |
| Number of measures solely taken by automated means | 16,747,079,313        |

## Content Moderation following a User Report

Microsoft provides the information below in response to DSA Articles 15(1)(b)-(c) and 41(2)(a)-(b).

Microsoft Advertising provides access for users to report ads they believe violate compliance associated with a particular harm using the Microsoft Advertising Report a Concern form located at <https://about.ads.microsoft.com/en/forms/policies/report-spam-form>

- At the bottom of each page on Bing, users can find a link entitled “Feedback;” this link directs to a text box where users can share their views on search results, generated responses, or other aspects of the services including Microsoft Advertising ads. User feedback is triaged in accordance with Microsoft internal guidelines and used as insights for product development. Other publishers may have similar functionalities.
- When users have specific concerns about ads they see, they may also report concerns through the Microsoft Advertising Report a Concern tool, made independently accessible in addition to the above Feedback method and Bing Report a Concern form. All these methods escalate concerns to a team of reviewers for compliance assessment. Users can report:
  - Malvertising



- Disallowed Content
- Relevancy Concerns
- Users can report relevance, content or other issues including:
  - Disallowed Content: Bullying or harassment, Graphic violence and human gore, Threat or exposure of private sexual images, Child sexual abuse, Hate speech, Suicide or self-harm, offering or requesting sex, terrorism or violent extremism, coordinating to physically harm, threats or praise of violence, trafficking, fraud, phishing or scams, viruses, spyware or malware, infringement and other concerns not listed.
  - Relevance: ad is not relevant, landing page is not relevant, ad copy does not make sense, display URL does not match the landing page or other relevance issues.
  - Page or site quality: high percentage of ads or links on the landing page, low value, sparse or limited content across the site, site redirection to an unrelated location or domain, other page or site quality issues or spam.
  - Personally identifiable information (PII): site asks for unexpected personal information or phishing.
  - Landing page navigation: site changes browser preferences without consent, site spawns multiple pop ups or prevents leaving the site, landing page does not load, product not available messaging or other landing page navigation issues.
- Users may also report ads they believe may infringe on intellectual property rights through the Intellectual Property Complaint form.
- Microsoft Advertising conducts a manual review of advertisements flagged to its customer support team and removes advertisements that violate its policies. Users will receive notification after action is taken.

#### Product Data: Microsoft Advertising

| Policy Violation                  | Number of notices received | Number of specific items of information included in the total number of notices | Median time to take action | Number of actions taken on the basis of the terms and conditions of the service |
|-----------------------------------|----------------------------|---------------------------------------------------------------------------------|----------------------------|---------------------------------------------------------------------------------|
| IP Infringements                  | 8,284                      | 8,284                                                                           | 0.0                        |                                                                                 |
| Pornography or Sexualized Content | 163,463                    | 163,463                                                                         | 0.0                        |                                                                                 |
| Scams and Fraud                   | 1,418                      | 1,418                                                                           | 0.0                        | 128                                                                             |

|                                |         |         |     |     |
|--------------------------------|---------|---------|-----|-----|
| Scope of Platform Services     | 23,153  | 23,153  | 0.0 | 107 |
| Unsafe and Prohibited Products | 454,435 | 454,435 | 0.0 |     |

## Content Moderation Appeals

Microsoft provides the information below in response to DSA Article 15(1)(d)

|                                                                     |                              |
|---------------------------------------------------------------------|------------------------------|
|                                                                     | <b>Microsoft Advertising</b> |
| Number of complaints submitted to the internal-complaints mechanism | 3,673,531                    |