

Introduction

The Microsoft Community (answers.microsoft.com) is an online platform where people can get help with their Microsoft products. It is a large online forum where Microsoft product users can ask questions, share their experiences, offer solutions to each other, and view Microsoft posted support content related to various Microsoft products.

Customers find it to be a valuable resource because it harnesses the collective knowledge and experiences of many expert users and Microsoft support. Instead of contacting Microsoft customer support directly for help, users can reuse a similar solution that has already been issued, getting to a faster resolution. Overall, the Microsoft Community is a helpful place for users to connect and assist each other with Microsoft product usage.

Contents

This DSA Transparency Report is responsive to the obligations under DSA Article 15(1), Article 24(1), and Article 42(2)-(3). This report provides information regarding the following topics as they pertain to the European Union:

- Content Moderation at Microsoft's Initiative
- Content Moderation & Automated Systems
- Content Moderation following a User Report
- Government Requests

Content Moderation at Microsoft's Initiative

Summary

Microsoft provides the information below in response to DSA Article 15(1)(c). This section reports data regarding content moderation Microsoft Community engaged in on its own initiative, absent a user report.

This section describes activities the Microsoft Community undertakes at its own initiative to address illegal content or information in violation of Microsoft Community's terms and conditions, which can be found by following this link: <https://answers.microsoft.com/en-us/page/tou>

Microsoft Community uses:

- Automatic detection and remediation for abuse or unacceptable content
- Human moderation where moderators can act reactively on receipt of a user's report (described later in the report) or a notice from the automatic tools (listed below) about concerning content, or proactively, when they notice any type of abuse while monitoring the forums

Content that violates Microsoft policies, or is unlawful, is blocked from sharing, edited, or is removed. The user profile in the Community might be banned, blocking them from creating new content, if they violate <https://answers.microsoft.com/en-us/page/faq#faqCodeConduct> . In the case of egregious violations such as sharing child sexual abuse material, the Microsoft account of the content owner may also be terminated in accordance with Microsoft's published policies.

Automated content moderation

Microsoft Community uses automated content detection on all text and image content when it is being submitted to Community posts.

Microsoft Community leverages automated content detection to detect known child sexual abuse material (CSAM) and terrorist content using both PhotoDNA and MD5 hashes.

PhotoDNA is a perceptual hashing technology, which is used to identify visually identical and near identical images and video. To do so, PhotoDNA creates a unique digital signature (known as a "hash") of an image or selected frames in a video and compares it against signatures ("hashes") of previously identified CSAM and terrorist content images or video. PhotoDNA is not facial recognition software and cannot be used to identify a person or object in an image.

For CSAM, PhotoDNA is used to detect content that has previously been determined to sexualize a child based on a classification scale widely adopted across industry. The service draws from an internal database of Microsoft-created hashes for CSAM content, as well as hashes provided by trusted third-party organizations.

For terrorist content, PhotoDNA is used to detect content that has previously been determined to be violent extremist and terrorism content consistent with Microsoft policies and which has been hashed for inclusion in an internal database of Microsoft-created hashes, as well as hashes provided by the Global Internet Forum to Counter Terrorism ("GIFCT").

Message-Digest Algorithm ("MD5") is a broadly used cryptographic hashing technology, which is provided as part of standard programming libraries and is used to determine if two items of content are identical. Microsoft currently deploys MD5 in conjunction with PhotoDNA to identify CSAM and terrorist content.

There are several automatic tools Microsoft uses to scan for abuse and unacceptable content:

- **Profanity filter** – the abusive word tracked in a Microsoft managed master file are replaced with the asterisks.
- **Blocked domains filter** – Domains Microsoft has blocked are replaced with an asterisk.
- **Blocked phone numbers list** – whenever a blocked phone number is posted, the post and profile become hidden from public view and go into spam review queue in the admin dashboard.
- **Image scanner** – when an automatic scanner detects a violation in an image, the post will be reported as abusive.
- **Machine learning (ML) spam tool** – whenever the ML spam tool detects spam content, the post and profile become hidden from public view and are added to the spam review queue in the admin dashboard. The service restricts the content owner's ability to edit or share that content while the content awaits human review.
- **Email address obfuscator** – all email addresses posted by users are automatically obfuscated in order to protect their privacy.

Human content moderation

Microsoft Community moderators are assigned to monitor the forums and act proactively on any type of abuse they notice. The table below shows the various human content moderation actions that occur on the community content that has been posted by site users.

Role	Qualification/Affiliation	Permissions/Empowerment level
Volunteer Abuse Moderator	Trusted and tenured external Community users who are not Microsoft employees	Thread/reply level abuse management Possibility to temporarily suspend user.
Moderator	Service Delivery Advocates who are Microsoft contract staff	Thread/reply level abuse management. Possibility to temporarily suspend user
Forum Owner	Microsoft product group representatives who are responsible for a particular product forum, may be either Microsoft employees or contract staff	Thread/reply level abuse management. Possibility to temporarily suspend user
Site Admin	Community Engagement team (CET) and Community Quality audit team, comprising Microsoft employees and contract staff	Thread/Profile level abuse management Possibility to block user

Content Moderation & Automated Systems

Microsoft provides the information below in response to DSA Articles 15(1)(e) and 42(2)(c).

Qualitative description of the automated means

The specifics of the technical implementation of PhotoDNA and MD5 hash matching are covered above.

The PhotoDNA distance threshold, which determines how “close” the Euclidian distance between two hashes needs to be in order to consider them a match, was selected based on the largest distance for which there were no false positives as determined during PhotoDNA’s initial validation, performed in partnership with the National Center for Missing and Exploited Children.

Given that MD5 is solely capable of identifying exact matches, there are no parameters set for its use.

Microsoft Community performs automated detection by generating PhotoDNA and MD5 hashes for the content to be scanned and comparing the resulting hash against the hashes of known bad content (as earlier described). If a match is found, the suspected CSAM or terrorist content is verified through human review, with the following exceptions:

- If additional content that is awaiting review has the same hash, only one such instance will be reviewed by a human reviewer and the result of that review will be applied to all content pending review with that same hash.
- If there have been at least ten earlier human reviews that confirmed CSAM content from a given content owner within the previous day, then all additional content from that same content owner

that matches a CSAM hash will be automatically actioned on as confirmed CSAM without additional human review.

If content is confirmed to be violative as described above, Microsoft Community will block sharing or remove the content and may terminate the account in accordance with Microsoft's published policies. The service also notifies the content owner of the enforcement action and, if the content appears to be CSAM, reports it to the U.S. National Center for Missing and Exploited Children.

Qualitative description of indicators of accuracy and possible rate of error of automated means.

There are two primary contributors to the effectiveness of hash matching systems that leverage PhotoDNA or MD5: the correctness of the hashing algorithm itself (the technology) and the management of the set of hashes against which content is matched (the content). The technology itself is content agnostic.

The correctness of the technology is ensured by validating that any updates do not alter its results when tested against a known set of inputs (i.e., we confirm that the hash calculations return the same result).

Hashing algorithms leverage databases of reference hashes of previously identified CSAM and terrorist content images or videos when assessing potentially violative content to order to find copies of the same content. Given that PhotoDNA and MD5 hash values are guaranteed to be stable, the precision and recall of these algorithms is, by definition, 100% for a set of previously identified violative hashes.

The following information relates to how the performance of the content (i.e. reference hashes) is tested and monitored. This is separate and distinct from the performance of the technology and is impacted by match distances used and the quality and consistency of human reviews both by Microsoft and by the third parties that generate hashes that Microsoft leverages.

Microsoft's reference hashes are updated and maintained in the following ways:

- Human review improves the accuracy of the underlying hashes by considering what the reviewer determined.
- In addition, in the case of PhotoDNA, hashes that are very similar are grouped to ensure they provide the same result.

As an indicator of accuracy, we have reversed 158 instances of Terrorist and CSAM enforcement actions. There is no direct measure of recall, though it should be noted that the previously identified hashes represent a small but unknown fraction of all violative content.

As noted previously, human review is not leveraged in cases where it would have no impact on the outcome of the review process (such as the prevention of multiple reviews of duplicate content) or would have no significant impact on the resulting enforcement actions (such as when the content owner's account has already been terminated for repeated confirmed violations). The measures of precision and recall of such actions would be equivalent to that of the underlying hash sets as described above.

Specification of the precise purposes to apply automated means.

Microsoft performs automated detection for child sexual abuse material (CSAM) because Microsoft is committed to protecting children from online harm. We do not allow the exploitation of, harm, or threat of harm to children on our services. This includes banning the use of our services to further child sexual exploitation and abuse (CSEA). When we become aware of content containing apparent child sexual exploitation and abuse, Microsoft reports the content to the National Center for Missing and Exploited Children (NCMEC).

CSEA is any content or activity that harms or threatens to harm a child through exploitation, trafficking, extortion, endangerment, or sexualization. This includes the sharing of visual media that contains sexual content that involves or sexualizes a child. CSEA also includes grooming, which is the inappropriate interaction with children by contacting, private messaging, or talking with a child to ask for or offer sex or sexual content, sharing content that is sexually suggestive, and planning to meet with a child for sexual encounters. A child is anyone under 18 years old.

Microsoft performs automated detection for terrorist content because we want to protect our users from illegal and harmful content and conduct. We also recognize that we have an important role to play in preventing terrorists and violent extremists from abusing online platforms. We do not allow content that praises or supports terrorists or violent extremists, helps them to recruit, or encourages or enables their activities. We look to the United Nations Security Council's Consolidated List to identify terrorists or terrorist groups. Violent extremists include people who embrace an ideology of violence or violent hatred towards another group.

In addressing terrorist and violent extremist content, we also work to ensure that people can use our services to talk about terrorism or violent extremism, share news or research about it, or express opposition to it.

There are several automatic tools that scan for abuse in the Community:

- **Profanity filter** – the abusive word is replaced with the asterisks.
- **Banned domains filter** – the blocked domain is replaced with the asterisks. The list is managed by the Site Admin
- **Banned phone numbers list** – whenever the blocked phone number is posted, the post and profile become hidden from public view and go into Spam review queue in the admin dashboard
- **Image scanner** – when an automatic scanner detects a violation in an image, the post will be reported as abusive.
- **ML (machine learning) spam tool** - whenever the ML detects spam behavior, the post and profile become hidden from public view and go into Spam review queue in the admin dashboard
- **Email obfuscator** – all emails posted by users are automatically obfuscated in order to protect their privacy.

Microsoft chooses not to perform human review on content with a hash that matches the hash of another content item that was already awaiting human review, because the matching hash value confirms that additional reviews of that same content provide no additional benefit.

Microsoft chooses not to perform human review of additional suspected CSAM content from a given content owner after at least ten content items from that owner have been confirmed by human review to be CSAM within the previous day because, based on published policy, the associated account would have been terminated after the first such confirmation and, thus, the marginal value of reviewing additional violations is limited with respect to informing enforcement actions.

In both cases, the elimination of these low-value reviews reduces the frequency with which human reviewers are exposed to harmful content that could negatively impact their wellness.

Safeguards applied to the use of automated means.

Microsoft leverages a broader human review process in conjunction with hash-based automated content detection. A matched image is directed to a human moderation queue (depending on the harm type of the matching reference hash), where one or more human reviewers may independently review the content to confirm that the content is violative before enforcement action is taken.

In cases where no human review occurs because content with the same hash was already in the process of being reviewed or the owning account had already been terminated as a result of at least ten other content items that had been reviewed by humans and confirmed to be CSAM, the impacted content owner can submit an appeal for the return of the content.


In the case of a user whose account is terminated following the detection of CSAM by automated means, the impacted user can submit an appeal to have the account returned. As an additional safeguard, Microsoft monitors the frequency with which such appeals result in the account suspension being reversed.

	Microsoft Community
Number of measures not taken by automated means	27

Content Moderation following a User Report

Microsoft provides the information below in response to DSA Articles 15(1)(b)-(c) and 41(2)(a)-(b).

1. All Community visitors see the **Report abuse** button on any post/profile.



2. It is necessary to be logged in to report abuse.
3. Once logged in, users may select option to report abuse anonymously
4. After the abuse report is submitted, users who reported abuse will receive email notification (unless they opted to stay anonymous)

5. Moderators see reported abuse flag and can act on it

[I want that on my desktop Automatic wallpaper change](#)
Last reply: October 2, 2020 | Created: October 2, 2020

I want that on my desktop Automatic wallpaper change after some interval in windows 10.

Reported Locked Top Content

1.58K views 42 helpful 2 replies

[blue outline on everything i click on](#)
Last reply: August 22, 2019 by | Created: January 24, 2013 by

We have Windows 8 now. I've been trying to figure out how to get rid of the blue box (outline) that shows on everything I click on, and when I type. It's so annoying and when I use some applications and games it is in the way. Can someone please help? thank you Janelle

Reported Locked Acceptable Content

597K views 4.02K helpful 120 replies

6. After the abuse report is reviewed by the moderator the user who submitted report will receive email notification about the action (unless they opted to stay anonymous)

Product Data: MS Community

Policy Violation	Number of notices received	Number of specific items of information included in the total number of notices	Median time to take action	Number of actions taken on the basis of the terms and conditions of the service
Illegal or Harmful Speech	10	10	0	
IP Infringements	19	19	0	
Nonconsensual Behaviour	4	4	0	
Pornography or Sexualized Content	8	8	0	1
Protection of Minors	224	224	0	2
Scams and Fraud	16	16	0	
Scope of Platform Services	300	300	0	24
Self Harm	2	2	0	

Government Requests

Microsoft provides the information below in response to DSA Article 15(1)(a).

Microsoft Community received 0 (zero) orders from government authorities to act against alleged illegal content during the reporting period.

Member State	Total Government Requests	Request Granted
N/A	0	0