

Digital Services Act

Transparency Report for Designer

February 2025

Introduction

This report will cover the following products and services which Microsoft Designer integrates with to provide an end-user experience for image and design creation. This includes the Designer Web & Mobile Apps and Integrations. Information related to M365 Copilot, into which Microsoft Designer integrates, has not been covered in this report.

This DSA Transparency Report is responsive to the obligations under DSA Article 15(1), Article 24(1), and Article 42(2)-(3). This Report provides information regarding the following topics as they pertain to the European Union:

- Content Moderation at Microsoft's Initiative
- Content Moderation & Automated Systems
- Content Moderation following a User Report
- Content Moderation Appeals
- Government Requests

Content Moderation at Microsoft's Initiative

Microsoft provides the information below in response to DSA Article 15(1)(c). This section describes activities Microsoft Designer consumer experiences (hereinafter "Designer") undertake at Microsoft's own initiative to address illegal content or information in violation of Designer's terms and conditions.


Automated Content Detection – Image Upload:

For Generative AI features in Designer, Designer leverages automated content detection for Child Sexual Exploitation and Abuse (CSEA). When Designer flags content through automated detection, the service may restrict the content owner's ability to edit or share that content while the content awaits human review.

Violating content is blocked from sharing or is removed, and, in the case of egregious violations such as sharing child sexual abuse material, the Microsoft account of the content owner may also be terminated in accordance with Microsoft's published policies.

Use of AI-Based Classifiers:

When using Designer's Generative AI features, a user may submit text or image queries as inputs to Microsoft's AI models, which then generate text and/or an image. These are known as "prompts" and "responses". Designer uses classifiers (machine learning models that help sort data into labeled classes or categories) and content filters on user queries and prompts to help mitigate harm or prevent misuse. Examples include requests for information that could potentially lead to users being unexpectedly exposed to self-harm, violence, graphic content, hateful content, or misleading information. Flags for these classifiers may lead to mitigations, such as not returning generated content to the user.



User Terms Enforcement:

Repeated attempted violations of the Designer Terms of Use may result in temporary or permanent suspension from Designer and other Services (as defined in the Microsoft Services Agreement).

Please note that Designer generally does not conduct "content moderation" as defined in the Digital Services Act, as content generated by Designer, even when requested by users, is not content "provided by recipients of the service." User queries and user prompts can trigger internal safety systems designed to ensure the services work as intended, including by preventing Designer from generating content that could violate Microsoft policies or otherwise expose users to potentially harmful content. These activities do not qualify as "content moderation" as defined in Article 3(t) of the Digital Services Act since any outputs generated by Designer (such as text or visual responses) are not content "provided by recipients of the service." Nevertheless, we have provided additional descriptions of how these systems operate in the interests of transparency.

Further Detail About Content Moderation at Microsoft's Initiative

In addition to the above, more detailed information about Designer own-initiative content moderation is below.

PhotoDNA:

Designer leverages automated content detection to detect known child sexual abuse material (CSAM) content using both PhotoDNA and MD5 hashes.

PhotoDNA is a perceptual hashing technology, which is used to identify visually identical and near identical images and video. To do so, PhotoDNA creates a unique digital signature (known as a “hash”) of an image or selected frames in a video and compares it against signatures (“hashes”) of previously identified CSAM content images or video. PhotoDNA is not facial recognition software and cannot be used to identify a person or object in an image.

For CSAM, PhotoDNA is used to detect content that has previously been determined to sexualize a child based on a classification scale widely adopted across industry. The service draws from an internal database of Microsoft-created hashes for CSAM content, as well as hashes provided by trusted third-party organizations.

Message-Digest Algorithm (“MD5”) is a broadly used cryptographic hashing technology, which is provided as part of standard programming libraries and is used to determine if two items of content are identical. Microsoft currently deploys MD5 in conjunction with PhotoDNA to identify CSAM content.

Use of AI-Based Classifiers:

When using Designer's Generative AI features, a user may submit text or image queries as inputs to Microsoft's AI models, which then generate text and/or an image. These are known as “prompts” and “responses”. Designer uses classifiers (machine learning models that help sort data into labeled classes or categories) and content filters on user queries and prompts to help mitigate harm or prevent misuse. Examples include requests for information that could potentially lead to users being unexpectedly exposed to self-harm, violence, graphic content, hateful content, or misleading information. Flags for these classifiers may lead to mitigations, such as not returning generated content to the user.

User Terms Enforcement:

Repeated attempted violations of the Designer Terms of Use may result in temporary or permanent suspension from Designer and other Services (as defined in the Microsoft Services Agreement).

Product Data, Designer:

Policy Violation	Number of measures taken at the provider's own initiative	Number of measures taken after detection with solely automated means	Visibility restriction Removal	Visibility restriction Disable
Protection of Minors	76	76	73	3

Policy Violation	Designer
Number of measures not taken by automated means	65
Number of measures solely taken by automated means	11

As noted above, Designer generally does not conduct “content moderation” as defined in the Digital Services Act, as content generated by Designer, even when requested by users, is not content “provided by recipients of the service.” Designer uses classifiers and content filters on user queries and prompts to help mitigate harm or prevent misuse. These classifiers and content filters can span different areas that do not necessarily align to one specific type of legal or policy violation and may be based on patterns or other signals of potential misuse. In addition, beyond the “Protection of Minors” category, due to limitations on the relevant technical tracking system, Designer cannot provide precise numbers regarding the number of users who were temporarily or permanently restricted, or who filed appeals, during this period.

However, we can estimate that for the period from February – December 2024, there were approximately 3,560 instances in which the recipient account's access to Designer was limited as a result of violations of the Designer Terms of Use. These actions were taken as a result of users attempting to bypass Microsoft safety systems. We estimate that Designer received 1,750 appeals of these actions.

Content Moderation & Automated Systems

Microsoft provides the information below in response to DSA Articles 15(1)(e) and 42(2)(c).

In addition to the above information about PhotoDNA and AI-Based Classifiers, more detailed information about PhotoDNA is provided below:

PhotoDNA:

The specifics of the technical implementation of PhotoDNA and MD5 hash matching are covered above.

The PhotoDNA distance threshold, which determines how “close” the Euclidian distance between two hashes needs to be in order to consider them a match, was selected based on the largest distance for which there were no false positives as determined during PhotoDNA’s initial validation, performed in partnership with the National Center for Missing and Exploited Children.

Given that MD5 is solely capable of identifying exact matches, there are no parameters set for its use.

Designer performs automated detection by generating PhotoDNA and MD5 hashes for the content to be scanned and comparing the resulting hash against the hashes of known bad content (as earlier described). If a match is found, the suspected CSAM or terrorist content is verified through human review, with the following exceptions:

If additional content that is awaiting review has the same hash, only one such instance will be reviewed by a human reviewer and the result of that review will be applied to all content pending review with that same hash.

If there have been at least ten earlier human reviews that confirmed CSAM content from a given content owner within the previous day, then all additional content from that same content owner that matches a CSAM hash will be automatically actioned on as confirmed CSAM without additional human review.

If content is confirmed to be violative as described above, Designer will block sharing or remove the content and may terminate the account in accordance with Microsoft’s published policies. The service also notifies the content owner of the enforcement action and, if the content appears to be CSAM, reports it to the U.S. National Center for Missing and Exploited Children.

PhotoDNA:

There are two primary contributors to the effectiveness of hash matching systems that leverage PhotoDNA or MD5: the correctness of the hashing algorithm itself (the technology) and the management of the set of hashes against which content is matched (the content). The technology itself is content agnostic.

The correctness of the technology is ensured by validating that any updates do not alter its results when tested against a known set of inputs (i.e., we confirm that the hash calculations return the same result).

Hashing algorithms leverage databases of reference hashes of previously identified CSAM and terrorist content images or videos when assessing potentially violative content to order to find copies of the same content. Given that PhotoDNA and MD5 hash values are guaranteed to be stable, the precision and recall of these algorithms is, by definition, 100% for a set of previously identified violative hashes.

The following information relates to how the performance of the content (i.e. reference hashes) is tested and monitored. This is separate and distinct from the performance of the technology and is impacted by match distances used and the quality and consistency of human reviews both by Microsoft and by the third parties that generate hashes that Microsoft leverages.

Microsoft's reference hashes are updated and maintained in the following ways:

- Human review improves the accuracy of the underlying hashes by considering what the reviewer determined.
- In addition, in the case of PhotoDNA, hashes that are very similar are grouped to ensure they provide the same result.

As an indicator of accuracy, we have reversed 158 instances of Terrorist and CSAM enforcement actions. There is no direct measure of recall, though it should be noted that the previously identified hashes represent a small but unknown fraction of all violative content.

As noted previously, human review is not leveraged in cases where it would have no impact on the outcome of the review process (such as the prevention of multiple reviews of duplicate content) or would have no significant impact on the resulting enforcement actions (such as when the content owner's account has already been terminated for repeated confirmed violations). The measures of precision and recall of such actions would be equivalent to that of the underlying hash sets as described above.

PhotoDNA:

Microsoft performs automated detection for child sexual abuse material (CSAM) because Microsoft is committed to protecting children from online harm. We do not allow the exploitation of, harm, or threat of harm to children on our services. This includes banning the use of our services to further child sexual exploitation and abuse (CSEA). When we

become aware of content containing apparent child sexual exploitation and abuse, Microsoft reports the content to the National Center for Missing and Exploited Children (NCMEC).

CSEA is any content or activity that harms or threatens to harm a child through exploitation, trafficking, extortion, endangerment, or sexualization. This includes the sharing of visual media that contains sexual content that involves or sexualizes a child. CSEA also includes grooming, which is the inappropriate interaction with children by contacting, private messaging, or talking with a child to ask for or offer sex or sexual content, sharing content that is sexually suggestive, and planning to meet with a child for sexual encounters. A child is anyone under 18 years old.

Microsoft chooses not to perform human review on content with a hash that matches the hash of another content item that was already awaiting human review, because the matching hash value confirms that additional reviews of that same content provide no additional benefit.

Microsoft chooses not to perform human review of additional suspected CSAM content from a given content owner after at least ten content items from that owner have been confirmed by human review to be CSAM within the previous day because, based on published policy, the associated account would have been terminated after the first such confirmation and, thus, the marginal value of reviewing additional violations is limited with respect to informing enforcement actions.

In both cases, the elimination of these low-value reviews reduces the frequency with which human reviewers are exposed to harmful content that could negatively impact their wellness.

AI-Based Classifiers:

The precise purpose to apply automated means is for Microsoft's Generative AI services is to detect harmful user-generated and AI-generated content in applications and services. Microsoft also remains committed to its Responsible AI principles and employs its classifiers consistent with this approach.

PhotoDNA:

Microsoft leverages a broader human review process in conjunction with hash-based automated content detection. A matched image is directed to a human moderation queue

(depending on the harm type of the matching reference hash), where one or more human reviewers may independently review the content to confirm that the content is violative before enforcement action is taken.

In cases where no human review occurs because content with the same hash was already in the process of being reviewed or the owning account had already been terminated as a result of at least ten other content items that had been reviewed by humans and confirmed to be CSAM, the impacted content owner can submit an appeal for the return of the content.

In the case of a user whose account is terminated following the detection of CSAM by automated means, the impacted user can submit an appeal to have the account returned. As an additional safeguard, Microsoft monitors the frequency with which such appeals result in the account suspension being reversed.

Content Moderation following a User Report

Microsoft provides the information below in response to DSA Articles 15(1)(b)-(c) and 41(2)(a)-(b).

Designer does not have an in-product “report a concern” feature because Designer does not currently include sharing or collaboration features. This section is not applicable to Designer.

Content Moderation Appeals

Microsoft provides the information below in response to DSA Article 15(1)(d)

Please see “Further Detail About Content Moderation at Microsoft’s Initiative” section, above.

Government Requests

Microsoft provides the information below in response to DSA Article 15(1)(a).

Designer did not receive any applicable government requests during the reporting period.