

Digital Services Act

Transparency Report for Microsoft Forms

February 2025

Introduction

This report covers Microsoft Forms. Microsoft Forms is a web-based application where customers design and analyze forms, polls, quizzes.

This DSA Transparency Report is responsive to the obligations under DSA Article 15(1), Article 24(1), and Article 42(2)-(3). This Report provides information regarding the following topics as they pertain to the European Union:

- Content Moderation at Microsoft's Initiative
- Content Moderation & Automated Systems
- Content Moderation following a User Report
- Content Moderation Appeals
- Government Requests

Content Moderation at Microsoft's Initiative

Microsoft provides the information below in response to DSA Article 15(1)(c). This section reports data regarding content moderation Microsoft Forms engaged in on its own initiative, absent a user report.

This section describes activities Microsoft Forms undertakes at its own initiative to address illegal content or information in violation of Microsoft Forms terms and conditions, which can be found [Microsoft Services Agreement](#).

When Microsoft Forms flags content through automated detection, the service may restrict the content owner's ability to edit or share that content while the content awaits human review.

Violating content is blocked from sharing or is removed, and, in the case of egregious violations such as sharing child sexual abuse material, the Microsoft account of the content owner may also be terminated in accordance with Microsoft's published policies.

Microsoft Forms uses automated content detection on all image content when it is uploaded or shared by consumer accounts.

Microsoft Forms leverages automated content detection to detect known child sexual abuse material (CSAM) and terrorist content using both PhotoDNA and MD5 hashes. PhotoDNA is a perceptual hashing technology, which is used to identify visually identical and near identical images and video. To do so, PhotoDNA creates a unique digital signature (known as a “hash”) of an image or selected frames in a video and compares it against signatures (“hashes”) of previously identified CSAM and terrorist content images or video. PhotoDNA is not facial recognition software and cannot be used to identify a person or object in an image.

For CSAM, PhotoDNA is used to detect content that has previously been determined to sexualize a child based on a classification scale widely adopted across industry. The service draws from an internal database of Microsoft-created hashes for CSAM content, as well as hashes provided by trusted third-party organizations.

For terrorist content, PhotoDNA is used to detect content that has previously been determined to be violent extremist and terrorism content consistent with Microsoft policies and which has been hashed for inclusion in an internal database of Microsoft-created hashes, as well as hashes provided by the Global Internet Forum to Counter Terrorism (“GIFCT”).

Message-Digest Algorithm (“MD5”) is a broadly used cryptographic hashing technology, which is provided as part of standard programming libraries and is used to determine if two items of content are identical. Microsoft currently deploys MD5 in conjunction with PhotoDNA to identify CSAM and terrorist content.

The table below reports information regarding the number of pieces of EU-content Microsoft Forms actioned during the reporting period absent a user report, organized by policy violation. Microsoft Forms assigns each piece of content a single policy violation. For each category of policy violation, Microsoft Forms reports the number of pieces of content actioned and which actions were taken.

Policy Violation	Number of measures taken at the provider's own initiative	Number of measures taken after detection with solely automated means	Visibility restriction Removal	Account restriction Suspension
Protection of Minors	10	10	8	2

Content Moderation & Automated Systems

Microsoft provides the information below in response to DSA Articles 15(1)(e) and 42(2)(c).

The specifics of the technical implementation of PhotoDNA and MD5 hash matching are covered above.

The PhotoDNA distance threshold, which determines how “close” the Euclidian distance between two hashes needs to be in order to consider them a match, was selected based on the largest distance for which there were no false positives as determined during PhotoDNA’s initial validation, performed in partnership with the National Center for Missing and Exploited Children.

Given that MD5 is solely capable of identifying exact matches, there are no parameters set for its use.

Microsoft Forms performs automated detection by generating PhotoDNA and MD5 hashes for the content to be scanned and comparing the resulting hash against the hashes of known bad content (as earlier described). If a match is found, the suspected CSAM or terrorist content is verified through human review, with the following exceptions:

- If additional content that is awaiting review has the same hash, only one such instance will be reviewed by a human reviewer and the result of that review will be applied to all content pending review with that same hash.
- If there have been at least ten earlier human reviews that confirmed CSAM content from a given content owner within the previous day, then all additional content from that same content owner that matches a CSAM hash will be automatically actioned on as confirmed CSAM without additional human review.

If content is confirmed to be violative as described above, Microsoft Forms will block sharing or remove the content and may terminate the account in accordance with Microsoft’s published policies. The service also notifies the content owner of the enforcement action and, if the content appears to be CSAM, reports it to the U.S. National Center for Missing and Exploited Children.

There are two primary contributors to the effectiveness of hash matching systems that leverage PhotoDNA or MD5: the correctness of the hashing algorithm itself (the technology) and the management of the set of hashes against which content is matched (the content). The technology itself is content agnostic.

The correctness of the technology is ensured by validating that any updates do not alter its results when tested against a known set of inputs (i.e., we confirm that the hash calculations return the same result).

Hashing algorithms leverage databases of reference hashes of previously identified CSAM and terrorist content images or videos when assessing potentially violative content to order to find copies of the same content. Given that PhotoDNA and MD5 hash values are guaranteed to be stable, the precision and recall of these algorithms is, by definition, 100% for a set of previously identified violative hashes.

The following information relates to how the performance of the content (i.e. reference hashes) is tested and monitored. This is separate and distinct from the performance of the technology and is impacted by match distances used and the quality and consistency of human reviews both by Microsoft and by the third parties that generate hashes that Microsoft leverages.

Microsoft's reference hashes are updated and maintained in the following ways:

- Human review improves the accuracy of the underlying hashes by considering what the reviewer determines.
- In addition, in the case of PhotoDNA, hashes that are very similar are grouped to ensure they provide the same result.

As an indicator of accuracy, we have reversed 158 instances of Terrorist and CSAM enforcement actions. There is no direct measure of recall, though it should be noted that the previously identified hashes represent a small but unknown fraction of all violative content.

As noted previously, human review is not leveraged in cases where it would have no impact on the outcome of the review process (such as the prevention of multiple reviews of duplicate content) or would have no significant impact on the resulting enforcement actions (such as when the content owner's account has already been terminated for repeated confirmed violations). The measures of precision and recall of such actions would be equivalent to that of the underlying hash sets as described above.

The precise purpose of these automated means is to prevent the dissemination of known CSEA and TVEC on Microsoft Forms.

Microsoft performs automated detection for child sexual abuse material (CSAM) because Microsoft is committed to protecting children from online harm. We do not allow the exploitation of, harm, or threat of harm to children on our services. This includes banning

the use of our services to further child sexual exploitation and abuse (CSEA). When we become aware of content containing apparent child sexual exploitation and abuse, Microsoft reports the content to the National Center for Missing and Exploited Children (NCMEC).

CSEA is any content or activity that harms or threatens to harm a child through exploitation, trafficking, extortion, endangerment, or sexualization. This includes the sharing of visual media that contains sexual content that involves or sexualizes a child. CSEA also includes grooming, which is the inappropriate interaction with children by contacting, private messaging, or talking with a child to ask for or offer sex or sexual content, sharing content that is sexually suggestive, and planning to meet with a child for sexual encounters. A child is anyone under 18 years old.

Microsoft performs automated detection for terrorist content because we want to protect our users from illegal and harmful content and conduct. We also recognize that we have an important role to play in preventing terrorists and violent extremists from abusing online platforms. We do not allow content that praises or supports terrorists or violent extremists, helps them to recruit, or encourages or enables their activities. We look to the United Nations Security Council's Consolidated List to identify terrorists or terrorist groups. Violent extremists include people who embrace an ideology of violence or violent hatred towards another group.

In addressing terrorist and violent extremist content, we also work to ensure that people can use our services to talk about terrorism or violent extremism, share news or research about it, or express opposition to it.

Microsoft chooses not to perform human review on content with a hash that matches the hash of another content item that was already awaiting human review, because the matching hash value confirms that additional reviews of that same content provide no additional benefit.

Microsoft chooses not to perform human review of additional suspected CSAM content from a given content owner after at least ten content items from that owner have been confirmed by human review to be CSAM within the previous day because, based on published policy, the associated account would have been terminated after the first such confirmation and, thus, the marginal value of reviewing additional violations is limited with respect to informing enforcement actions.

In both cases, the elimination of these low-value reviews reduces the frequency with which human reviewers are exposed to harmful content that could negatively impact their wellness.

Microsoft leverages a broader human review process in conjunction with hash-based automated content detection. A matched image is directed to a human moderation queue (depending on the harm type of the matching reference hash), where one or more human reviewers may independently review the content to confirm that the content is violative before enforcement action is taken.

In cases where no human review occurs because content with the same hash was already in the process of being reviewed or the owning account had already been terminated as a result of at least ten other content items that had been reviewed by humans and confirmed to be CSAM, the impacted content owner can submit an appeal for the return of the content.


In the case of a user whose account is terminated following the detection of CSAM by automated means, the impacted user can submit an appeal to have the account returned. As an additional safeguard, Microsoft monitors the frequency with which such appeals result in the account suspension being reversed.

The table below reports how many pieces of EU-content were automatically processed and how many were manually processed, broken down by how they were detected.

	Microsoft Forms
Number of measures not taken by automated means	15

Content Moderation following a User Report

Microsoft provides the information below in response to DSA Articles 15(1)(b)-(c) and 41(2)(a)-(b).



Untitled form

1. Question

Enter your answer

Submit

Never give out your password. [Report abuse](#)

Report Abuse

What violation are you reporting?

- ☐ Phishing
- ☐ Threats, cyberbullying, harassment
- ☐ Child endangerment or exploitation
- ☐ Offensive content
- ☐ Virus, spyware, malware
- ☐ Copyright infringement
- ☐ Other

Comments

Please provide additional details about the abuse or offensive behavior you're reporting.

Please confirm the accuracy of the information you have provided. If you wish to receive updates regarding this report, you can choose to opt in using the options below.

Would you like to provide us with your contact info for any follow up questions and receive the results of our review?

- ☒ Yes
- ☐ No

Name *

Enter your full name

Email address *

Enter your email address

Submit

Report Abuse

What violation are you reporting?

☐ Phishing

☐ Threats, cyberbullying, harassment

☐ Child endangerment or exploitation

☐ Offensive content

☐ Virus, spyware, malware

☐ Copyright infringement

☐ Other

Comments

Please provide additional details about the abuse or offensive behavior you're reporting.

Please confirm the accuracy of the information you have provided. If you wish to receive updates regarding this report, you can choose to opt in using the options below.

Would you like to provide us with your contact info for any follow up questions and receive the results of our review?

☐ Yes

☒ No

Submit

The user-initiated report can be performed through the Report a Concern (RaC) page. When users encounter a form that they want to report, they can click on the “report abuse” button to access the RaC page. On this page, users can select a category for their report and may also provide additional comments. Users can opt-in to report with their contact information or report anonymously.

The table below reports information regarding the number of EU reports Microsoft Forms received during the reporting period by user-selected report reason. For each report reason, Microsoft Forms provides the number of reports received, the pieces of content underlying those reports, and the number of reports where Microsoft Forms determined the content violated its policies.

Policy Violation	Number of notices received	Number of specific items of information included in the total number of notices	Median time to take action	Number of actions taken on the basis of the terms and conditions of the service
Nonconsensual Behaviour	571	571		
Pornography or Sexualized Content			10.5	2
Protection of Minors	54	54		
Scams and Fraud	242	242	3.0	3
Scope of Platform Services	449	449		

Content Moderation Appeals

Microsoft provides the information below in response to DSA Article 15(1)(d)

Microsoft Forms offers users the option to appeal or request a re-review of an item that has been blocked from sharing. Users can initiate the appeal process via Microsoft customer support channel. If an item has been permanently deleted due to its classification as Child Sexual Exploitation and Abuse (CSEA) material, it cannot be appealed through the inline product approach. In such cases, the user must follow the Microsoft Services Agreement (MSA) appeal process to regain account access; however, the content in question will remain permanently deleted.

Government Requests

Microsoft Forms provides the information below in response to DSA Article 15(1)(a)

Microsoft Forms carefully evaluates all government requests for content removal. For government demands, Microsoft Forms integrate with central service to ensure any actions taken are narrow, specific, submitted in writing, and based on valid legal orders.

During the reporting period, Microsoft Forms did not have any government requests from EU Member State government authorities.