

Digital Services Act

Transparency Report for OneDrive

February 2025

Introduction

This report will cover the following products and services as all of these services can utilize OneDrive storage and are aggregated in the database.

- Excel Online
- Lists
- Loop
- OneDrive
- OneNote
- PowerPoint Online
- Sticky Notes
- ToDo
- Visio
- Word

These are primarily consumer productivity products and are used by customers to increase productivity in their day-to-day lives.

This DSA Transparency Report is responsive to the obligations under DSA Article 15(1), Article 24(1), and Article 42(2)-(3). This Report provides information regarding the following topics as they pertain to the European Union:

- Content Moderation at Microsoft's Initiative
- Content Moderation & Automated Systems
- Content Moderation following a User Report
- Content Moderation Appeals
- Government Requests

Content Moderation at Microsoft's Initiative

OneDrive provides the information below in response to DSA Article 15(1)(c). This section reports data regarding content moderation OneDrive engaged in on its own initiative, absent a user report.

The content moderation that OneDrive performs at its own initiative is automated content detection for Child Sexual Exploitation and Abuse (CSEA) and Terrorist and Violent Extremist Content (TVEC). That automated detection is detailed in answers (2), (4), (5), (6), and (7). If content is deemed against TVEC policy, it will be restricted from visibility by users other than the content owner (block sharing capabilities). If content is deemed against CSEA policy, it will be permanently removed from the service and reported to the proper authorities as legally required.

The automated detection methods used by OneDrive are hash matching using PhotoDNA and MD5 hashes for the purpose of detecting CSEA and TVEC. Microsoft reviews content when shared from OneDrive for compliance with our laws and policies, and in some jurisdictions, also reviews content stored in a users' OneDrive.

Engineers and Product Managers:

The Microsoft Digital Safety Standard and its associated program is the basis for informing engineers and product managers of their Digital Safety commitments. Including:

- The Microsoft Digital Safety Standard (MDSS): A high-level document that sets out objectives and basic requirements for safety systems and processes.
- Implementation Guidance: Granular requirements for specific types of services to enable compliance with a wide variety of laws requiring safety systems and processes. Facilitates engineering and operational activities.
- Content Policy Framework: Common methodology for development of content policies, which are foundational documentation for new regulations requiring transparency of moderation practices and consistency in application.
- Service Classification: Enables law-to-service mapping,
- Additional Requirements: Catalogs variable legal obligations across enacted laws.
- Governance: Details the oversight and responsibilities for services' implementation of MDSS objectives.

The above information is shared through various training opportunities, including:

- Introduction to Digital Safety
- Regular MDSS trainings (includes country-specific requirements)

- Monthly topical learning sessions, e.g., technology and tools for digital safety at Microsoft; specific content harms and risks.
- Role-based job training, e.g., content review analysts; compliance
- Digital Safety Program Owners (product group members accountable for a service's digital safety posture) and Digital Safety Accountable Leaders (executives accountable for the digital safety posture of a portfolio of products): knowledge sharing and support, and 1:1 mentoring

Human Reviewers:

A content policy codifies the definition of a given harm. An operations team then takes that as input and produces reviewer guidelines, that provide concrete prescriptive guidance on how reviewers should assess content along with examples. The team that manages the reviewers that are assigned to that harm is then tasked with operationalizing those guidelines.

Consistency and correctness are then maintained through a collection of automated and manual processes including:

- Monitoring of agreement rates for content that undergoes multiple independent reviews, either in parallel, as occurs by policy in some cases involving CSAM, or in cases where the same/similar content is encountered multiple times and is, therefore, sent for review multiple times.
- Processes that are run by review teams to revisit nuanced cases to ensure there is consensus on how they should be handled and that the appropriate clarifications are made to reviewer guidelines as a result.
- There are monthly operational reviews that will surface any issues found by either of the above that cannot be quickly resolved to higher-level management.

It is also important to acknowledge the assistance that is provided to reviewers due to the mentally and emotionally taxing nature of their work. Where possible, we leverage tools with built-in wellness functionality (e.g. allowing images to be converted into black and white, to be blurred, to have faces blocked, to mute sound in videos, etc). We also ensure that reviewers have access to mental well-being services.

The table below reports information regarding the number of pieces of EU-content OneDrive actioned during the reporting period absent a user report, organized by policy violation. OneDrive assigns each piece of content a single policy violation. For each category of policy violation, OneDrive reports the number of pieces of content actioned and which actions were taken.

OneDrive did not have any government-designated trusted flaggers for the EU for this time period, and therefore there are no reports coming from that source.

Product Data, respectively: Lists | Loop | OneDrive

Part 1 of 2

Policy Violation	Number of measures taken at the provider's own initiative			Number of measures taken after detection with solely automated means			Visibility restriction Removal		
Pornography or Sexualized Content		2	17		2	17			
Protection of Minors	8	509	11864	8	509	11864	5	443	10573
Risk for public security		10	106		10	106			
Violence			5			5			

Part 2 of 2

Policy Violation	Visibility restriction Disable			Visibility restriction Interaction restricted			Account restriction Suspension		
Pornography or Sexualized Content		2	17						
Protection of Minors		32	709				3	34	582
Risk for public security		1	7		9	99			
Violence			5						

Content Moderation & Automated Systems

OneDrive provides the information below in response to DSA Articles 15(1)(e) and 42(2)(c).

The specifics of the technical implementation of PhotoDNA and MD5 hash matching is covered in answer (2).

The in-scope service generates a hash for the content to be scanned and submits to a central service for matching. If a match is detected, the in-scope service is informed of the match and responds by sending the content and its associated metadata to the central service.

The central service orchestrates any and all human reviews that are required by policy and, if the content is confirmed to be violative, the centralized service will communicate the need to take any associated content-level actions (e.g. blocking sharing or deleting the content) back to the in-scope service. If a violation is confirmed on OneDrive, it will be blocked from sharing on the service and may be permanently removed from the service.

The central service then performs any additional required processing including notification of the action to the content owner, account suspension (if warranted), reporting any CSEA discovered to NCMEC, and persisting the evidence to enable adjudicating appeals and responding to any resulting law enforcement data requests.

There are two primary contributors to the effectiveness of hash matching systems that leverage PhotoDNA or MD5: the correctness of the hashing algorithm itself (the technology) and the management of the set of hashes against which content is matched (the content). The technology itself is content agnostic.

The correctness of the hashing algorithm itself is ensured by validating that any updates do not alter its results on a known set of inputs (i.e., we confirm that the hash calculations return the same result, which is necessary to achieve 100% recall).

Hash matching leverages databases of reference hashes of known CSEA and Terrorism images or videos when assessing potentially violative content to find copies of the same content.

The following information relates to how the performance of the content (i.e. reference hashes) is tested and monitored. This is separate and distinct from the performance of the technology and is impacted by match distances used and the quality and consistency of human reviews both by Microsoft and by third parties that generate hashes that Microsoft leverages.

Reference hashes are updated and maintained in the following ways:

- Human review improves the accuracy of the underlying hashes by considering what the reviewer determines.
- In addition, in the case of PhotoDNA, hashes that are very similar are grouped to ensure they provide the same result.

As hash sharing entities provide new hashes, they are merged with the database to ensure consistent results.

Microsoft has no specific acceptance threshold for the rate at which matched content is determined to be CSEA or Terrorism content beyond our goal to improve the rate over time. Microsoft targets a rate of less than 5% for matched content that is determined to be benign, because historically, it has been able to achieve that rate without losing recall. Given the nature of hash matching and the inherent differences of both policies and training between Microsoft and our industry peers, we cannot expect 100% precision. As a result, our thresholds are informed by historical performance, and we carefully weigh any improvements in precision in either rate against the potential for loss of recall.

The precise purpose of these automated means is to prevent the dissemination of known CSEA and TVEC on OneDrive.

Microsoft leverages a broader human review process in conjunction with hash-based detection. A matched image is directed to a human moderation queue (depending on the harm type of the matching reference hash) where one or more moderators may independently review the content for quality control purposes. In order for a hash signature to be added to Microsoft's internal database of verified hashes, the associated image or video frame must have been reviewed by at least two human reviewers.


The table below reports how many pieces of EU-content were automatically processed and how many were manually processed, broken down by how they were detected.

	Lists	Loop	OneDrive
Number of measures not taken by automated means	8	552	12,520
Number of measures solely taken by automated means		5	725

Content Moderation following a User Report

OneDrive provides the information below in response to DSA Articles 15(1)(b)-(c) and 41(2)(a)-(b).

The user-initiated report can be performed through the Report a concern (RaC) dialog. When a user encounters a file in OneDrive that they want to report, they can click on the “...” menu and choose the Report a concern menu item.



In the RaC dialog, the user must choose an option provided and can optionally add comments with the report.

If a user wants to report anonymously, they can go to <https://www.microsoft.com/en-us/DigitalSafety/report-a-concern> to fill out a report as shown below.

Report a concern

We work diligently to protect our users from harmful content and conduct, but we can't do it alone. Your reports help us identify where online content or user conduct has violated our policies or terms for Microsoft services.

When you report content or conduct to Microsoft, we review it and act against it if it violates our [policies](#) or [terms](#). Often the easiest way to submit a report is within the product experience.

If you can't find an option to report digital safety abuse in the context in which you encountered it, select 'Continue' below to report it here.

If you have technical support concerns, please visit [Microsoft Support](#). If you have a privacy concern, you can [Report a Privacy concern](#).

If your report relates to deceptive generative AI election content please visit [Microsoft-2024 Elections](#).

Continue

Report a concern

Please provide as much information as possible about where you encountered the issue to help us address your concern. **An accurate link or web address is necessary** for us to complete a successful investigation.

Service or product *

OneDrive

Enter a link to the content you are reporting *

<https://1drv.ms/i/c/cc7f07bb5016cad2/EbCazEQxxq5PqpggNI>

How to get a OneDrive link

To get a link for the content you are reporting, please see the 'copy link' section of this article for instructions. Do not submit folder links because we can only process reports on single file links.

[How to share OneDrive files](#)

Go back

Continue

Report a concern

Please specify the type of concern and provide a detailed description to help us review the matter thoroughly.

Concern type *

Please select a concern type to continue

- Bullying or harassment
- Graphic violence and human gore
- Threat or exposure of private sexual images
- Child sexual abuse
- Hate speech
- Suicide or self-harm
- Offering or requesting sex
- Terrorism or violent extremism
- Coordinating to physically harm
- Threats or praise of violence
- Trafficking
- Spam
- Fraud, phishing, or scam
- Virus, spyware, or malware
- Infringement
- Other terms of use violation

Report a concern

Please specify the type of concern and provide a detailed description to help us review the matter thoroughly.

Concern type *

Terrorism or violent extremism

Description *

This is disturbing content

Go back

Continue

Report a concern

Please confirm the accuracy of the information you have provided. If you wish to receive updates regarding this report, you can choose to opt in using the options below.

Would you like to provide us with your contact info for any follow-up questions and receive the results of our review?

☐ Yes ☒ No

Microsoft is committed to helping our customers and to protecting your privacy. You may read the [Microsoft Privacy Statement](#) if you would like more information on how we protect customer information.

☒ I certify this report to be correct to the best of my knowledge.

Go back

Submit

The table below reports information regarding the number of EU reports OneDrive received during the reporting period by user-selected report reason. For each report reason, OneDrive provides the number of reports received, the pieces of content underlying those reports, and the number of reports where OneDrive determined the content violated its policies.

OneDrive did not have any government-designated trusted flaggers for the EU for this time period, and therefore there are no reports coming from that source.

Product Data, respectively: Lists | Loop | OneDrive

Part 1 of 2

Policy Violation	Number of notices received			Number of specific items of information included in the total number of notices		
Illegal or Harmful Speech		2	19		2	19
IP Infringements			18			128
Nonconsensual Behaviour			2			
Protection of Minors		1	7		1	6
Risk for public security			14			15
Scams and Fraud		40	1,089		40	1,088
Scope of Platform Services	7	65	1,237	7	65	1,225
Self Harm		1	46		1	46

Part 2 of 2

Policy Violation	Median time (hrs) to take action			Number of actions taken on the basis of the terms and conditions of the service		
IP Infringements			130.2			
Nonconsensual Behaviour			109.0			
Protection of Minors			129.7			
Risk for public security			772.0			
Scams and Fraud			48.8			11
Scope of Platform Services			18.2			
Violence			2.0			1

Content Moderation Appeals

OneDrive exposes the option to re-review an item that has been blocked from sharing. This is done inline in the OneDrive product. In this case, the customer retains access and edit rights to a file, so going through a traditional appeals process will not guarantee a review of the same content that is on the service.

If an item has been permanently deleted, it was because the item was confirmed as CSEA and cannot be appealed through the inline product approach. The user must go through the MSA appeal process to regain access to their account but that content will still remain permanently deleted from the OneDrive account. If the appeal is approved, the original content will be emailed back to the customer so they still retain a copy of their content. The numbers below reflect the MSA appeals process.

	Lists	Loop	OneDrive
Number of complaints submitted to the internal-complaints mechanism	1.0	4.0	627.0

Government Requests

OneDrive provides the information below in response to DSA Article 15(1)(a)

This section reports data on requests from Member State government authorities: (1) to remove content and (2) to provide user account information. OneDrive carefully considers all government requests for content removal and account information and works to mitigate any implications they may have on freedom of expression and human rights. For government demands, OneDrive employs safeguards to ensure any actions taken are narrow, specific, submitted in writing, and based on valid legal orders. Through its parent company, Microsoft, OneDrive also engages with broader civil society organizations on best practices related to government requests and participates in human rights impact assessments.

The table below reports information regarding the number of requests OneDrive received from Member State government authorities to remove content during the reporting period, organized by Member State and illegal content type. Government requests to remove content include requests reporting violations of our terms of service or violations of local law. OneDrive only received one government request during this reporting period.

Product Data: OneDrive

Member State	Number of orders to provide information	Median time to inform the authority of the receipt of the order to provide information	Median time to give effect to the order to provide information
Denmark	1	176	176
EU Internet Referral Unit	1	93	93
Germany	1	5	5