

Digital Services Act

Transparency Report for Microsoft Whiteboard

February 2025

Introduction

This report will cover the Microsoft Whiteboard product. Microsoft Whiteboard in Microsoft 365 is a free-form, digital new canvas that is used by consumers to brainstorm, collaborate and create diagrams. Whiteboard is a consumer productivity tool and lets team members collaborate in real time.

This DSA Transparency Report is responsive to the obligations under DSA Article 15(1), Article 24(1), and Article 42(2)-(3). This Report provides information regarding the following topics as they pertain to the European Union:

- Content Moderation at Microsoft's Initiative
- Content Moderation & Automated Systems
- Content Moderation following a User Report
- Content Moderation Appeals
- Government Requests

Content Moderation at Microsoft's Initiative

Microsoft provides the information below in response to DSA Article 15(1)(c).

(1) Summary of the content moderation engaged in at the provider's own initiative

This section describes activities Whiteboard undertakes at its own initiative for content moderation, absent a user report, to address illegal content or information in violation of Whiteboard's terms and conditions, which can be found [Microsoft Services Agreement](#).

Whiteboard with violating content is blocked or is removed, and in the case of egregious violations such as sharing child sexual abuse material, the Microsoft account of the content owner may also be terminated in accordance with Microsoft's published policies.

(2) Meaningful and comprehensible information regarding content moderation engaged in at the provider's own initiative

The content moderation that Whiteboard performs at its own initiative is automated content detection for Child Sexual Exploitation and Abuse (CSEA) and Terrorist and Violent Extremist Content (TVEC).

When Whiteboard flags content through automated detection, the service restricts the content owner's ability to edit or share that content while the content awaits human review. If content is deemed against TVEC policy, the whiteboard containing it will be restricted for all users, including the content owner, from viewing, editing or sharing, and access to the violating board is removed for all. If content is deemed against CSEA policy, it will be permanently removed from the service and reported to the proper authorities as legally required.

Whiteboard with violating content is blocked or is removed, and in the case of egregious violations such as sharing child sexual abuse material, the Microsoft account of the content owner may also be terminated in accordance with Microsoft's published policies.

Currently, Whiteboard is unable to filter EU-only data, but there is a plan to provide EU data in future reports. Therefore, the table below reports information regarding the number of pieces of Global-content Whiteboard actioned during the reporting period absent a user report, organized by policy violation. Whiteboard assigns each piece of content a single policy violation. For each category of policy violation, Whiteboard reports the number of pieces of content actioned and which actions were taken.

Policy Violation	Number of measures taken at the provider's own initiative	Number of measures taken after detection with solely automated means	Visibility restriction Removal	Account restriction Suspension
Protection of Minors	11	11	5	5

Content Moderation & Automated Systems

Microsoft provides the information below in response to DSA Articles 15(1)(e) and 42(2)(c).

(3) Qualitative description of Automated means

Whiteboard leverages Content Validation Service (CVS) to detect child sexual abuse material (CSAM) and Terrorist and Violent Extremist Content (TVEC).

Currently, Whiteboard scans for images uploaded by users in the boards for any violations. Additionally, Whiteboard scans the thumbnails of the boards as seen on the app home page. Whiteboard employs CVS for automated content detection of all images. CVS is a centralized Microsoft service that scans content to identify violations of Microsoft policies or Terms of Use (ToU).

For Image scanning using Content Validation Service (CVS), content items are compared to known CSAM PhotoDNA hashes from the National Center for Missing and Exploited Children (NCMEC), Microsoft, and the Tech Coalition to identify CSAM content.

(4) Qualitative description of indicators of accuracy and possible rate of error of automated means

PhotoDNA is a perceptual hashing technology, which is used to identify visually identical and near identical images and video. To do so, PhotoDNA creates a unique digital signature (known as a “hash”) of an image or selected frames in a video and compares it against signatures (“hashes”) of previously identified CSAM and terrorist content images or video. PhotoDNA is not facial recognition software and cannot be used to identify a person or object in an image.

The PhotoDNA distance threshold, which determines how close the Euclidean distance between two hashes must be to qualify as a match, was selected based on the largest distance for which no false positives were detected. This validation was conducted during PhotoDNA’s initial testing in partnership with NCMEC.

The correctness of the technology is ensured by validating that any updates do not alter its results when tested against a known set of inputs (i.e., we confirm that the hash calculations return the same result).

For CSAM, PhotoDNA is used to detect content that has previously been determined to sexualize a child based on a classification scale widely adopted across industry. The service draws from an internal database of Microsoft-created hashes for CSAM content, as

For terrorist content, PhotoDNA is used to detect content that has previously been determined to be violent extremist and terrorism content consistent with Microsoft policies and which has been hashed for inclusion in an internal database of Microsoft-created hashes, as well as hashes provided by the Global Internet Forum to Counter Terrorism (“GIFCT”).

(5) Specifications of the precise purposes to apply automated means

The precise purpose of these automated means is to prevent the dissemination of known CSEA and TVEC on Whiteboard.

Microsoft performs automated detection for child sexual abuse material (CSAM) because Microsoft is committed to protecting children from online harm. We do not allow the exploitation of, harm, or threat of harm to children on our services. This includes banning the use of our services to further child sexual exploitation and abuse (CSEA). When we become aware of content containing apparent child sexual exploitation and abuse, Microsoft reports the content to the National Center for Missing and Exploited Children (NCMEC).

Microsoft performs automated detection for terrorist content because we want to protect our users from illegal and harmful content and conduct. We also recognize that we have an important role to play in preventing terrorists and violent extremists from abusing online platforms. We do not allow content that praises or supports terrorists or violent extremists, helps them to recruit, or encourages or enables their activities. We look to the United Nations Security Council’s Consolidated List to identify terrorists or terrorist groups. Violent extremists include people who embrace an ideology of violence or violent hatred towards another group.

(6) Safeguards applied to use of automated means

Microsoft leverages a broader human review process in conjunction with hash-based automated content detection. A matched image is directed to a human moderation queue (depending on the harm type of the matching reference hash), where one or more human reviewers may independently review the content to confirm that the content is violative before enforcement action is taken.

Microsoft chooses not to perform human review on content with a hash that matches the hash of another content item that was already awaiting human review, because the matching hash value confirms that additional reviews of that same content provide no additional benefit.

Microsoft chooses not to perform human review of additional suspected CSAM content from a given content owner after at least ten content items from that owner have been confirmed by human review to be CSAM within the previous day because, based on published policy, the associated account would have been terminated after the first such confirmation and, thus, the marginal value of reviewing additional violations is limited with respect to informing enforcement actions.

In cases where no human review occurs because content with the same hash was already in the process of being reviewed or the owning account had already been terminated as a result of at least ten other content items that had been reviewed by humans and confirmed to be CSAM, the impacted content owner can submit an appeal for the return of the content.

In the case of a user whose account is terminated following the detection of CSAM by automated means, the impacted user can submit an appeal to have the account returned. As an additional safeguard, Microsoft monitors the frequency with which such appeals result in the account suspension being reversed.


The table below reports how many pieces of Global-content were automatically processed and how many were manually processed, broken down by how they were detected.

	Whiteboard
Number of measures not taken by automated means	9
Number of measures solely taken by automated means	1

Content Moderation following a User Report

Microsoft provides the information below in response to DSA Articles 15(1)(b)-(c) and 41(2)(a)-(b).

Users can report a concern on a Whiteboard file from the Whiteboard home page. When a user encounters a file in Whiteboard that they want to report, they can click on the “...” menu and choose the “Report” menu item as shown in the next set of screenshots.



In the Report concern dialog, the user must choose an option provided and can optionally add comments with the report.

The image displays two versions of a 'Report abuse' dialog box. The left version shows the main form with a title bar, a close button, and a question 'How does this break the code of conduct?'. Below this is a dropdown menu labeled 'Please select your concern...'. A text area follows with the prompt 'You can also include relevant details on the offensive content:' and a character count '1024 out of 1024 characters remaining'. At the bottom, there are links for 'Privacy and Cookies' and 'Copyright claims procedure', a checked checkbox for 'Report anonymously', and a 'Send report' button. The right version shows the same dialog but with the dropdown menu expanded, revealing a list of concerns: 'Bullying and Harassment', 'Graphic Violence and Human Gore', 'Non-consensual Intimate Imagery & Intimate Extortion', 'Child Sexual Exploitation and Abuse', 'Hate Speech', 'Suicide and Self-Injury', 'Sexual Solicitation' (which is highlighted), 'Terrorism and Violent Extremism', 'Coordination of Harm', 'Violent Threats, Incitement, and Glorification of Violence', and 'Trafficking'.

Users have the ability to report concerns anonymously by unchecking the selection box named “Report anonymously”, as shown in the next screenshot.

Report abuse

×

How does this break the code of conduct?

Terrorism and Violent Extremism

▼

This is disturbing content

997 out of 1024 characters remaining

[Privacy and Cookies](#)

[Copyright claims procedure](#)

☐ Report anonymously

Send report

The table below reports information regarding the number of Global reports Whiteboard received during the reporting period by user-selected report reason. For each report reason, Whiteboard provides the number of reports received, the pieces of content underlying those reports, and the number of reports where Whiteboard determined the content violated its policies.

Policy Violation	Number of notices received	Number of specific items of information included in the total number of notices	Median time to take action	Number of actions taken on the basis of the terms and conditions of the service
Illegal or Harmful Speech	97	97		
Nonconsensual Behaviour	102	102		
Pornography or Sexualized Content	3	3	0.5	2
Protection of Minors	23	23	4.0	
Risk for public security	10	10		
Scope of Platform Services			0.9	37
Self Harm	20	20		
Violence	705	705		

Content Moderation Appeals

Microsoft provides the information below in response to DSA Article 15(1)(d)

The impacted content owner can submit an appeal for the return of the content in cases described earlier, and in case of blocked accounts, the impacted user can submit an appeal to have the account returned.

The table below reports information regarding the number of complaints submitted to the internal complaint mechanisms of Microsoft.

Government Requests

Microsoft provides the information below in response to DSA Article 15(1)(a)

This section reports data on requests from Member State government authorities: (1) to remove content and (2) to provide user account information. Microsoft carefully considers all government requests for content removal and account information and works to mitigate any implications they may have on freedom of expression and human rights. For government demands, Microsoft employs safeguards to ensure any actions taken are narrow, specific, submitted in writing, and based on valid legal orders.

The table below reports information regarding the number of requests received from Member State government authorities to remove content during the reporting period, organized by Member State and illegal content type. Government requests to remove content include requests reporting violations of our terms of service or violations of local law.