



European Union Terrorist Content Online Transparency Report

February 29, 2024

The following report is published in accordance with Articles 7(2) and 7(3) of EU Regulation 2021/784, addressing the dissemination of terrorist content online (“TCO” or the “Regulation”). This report contains information for a period of seven months, from 1 January 2023 to 31 December 2023.

1. GENERAL OBSERVATIONS

General observations outlining the efforts undertaken by Meta to eliminate content violating our Terms and Conditions on our platform.

Meta Platforms Ireland Limited, which provides the Facebook social networking app for users in the EU, aims to create a safe and trusted platform, where people can feel free to express themselves. We are clear that we do not allow people to post content that threatens violence, promotes certain criminal activities, or includes praise and support for dangerous organizations or individuals. In order to achieve this, we take a multi-faceted approach to addressing potentially criminal or harmful activity that violates our policies on the Facebook platform.

A. We maintain a set of globally applicable [Community Standards](#) that define what is and isn't allowed on Facebook. These standards apply to content worldwide and are integral to protecting both expression and personal safety on Facebook.

Facebook is committed to equal treatment of all users. The [Facebook Community Standards](#) are applied consistently to everyone on the platform, irrespective of their location. We collaborate with global experts in technology, public safety, and human rights in developing these policies.

Our reviewers enforce these standards using comprehensive guidelines, which ensure that decisions are as consistent as possible and that users are treated equally and impartially.

Our Community Standards prohibit a wide range of objectionable or harmful content, including content that:

- Promotes violence
- Facilitates or promotes certain criminal activities
- Promotes dangerous organizations and individuals
- Incites violence
- Threatens the safety of others
- Is considered hate speech

The Community Standards are created by a global team with a wide array of backgrounds, including those who have dedicated their careers to issues like online safety, hate speech, and terrorism. This team regularly seeks input from outside experts and organizations to help balance the different perspectives that exist on free expression and safety, and to better understand the impacts of our policies on different communities globally.

Our policies under our Community Standards are regularly updated to keep pace with changes happening online and offline around the world. Updates to our policies are tracked in the change log so people can see how they have evolved (e.g., [Dangerous Organizations and Individuals change log](#)).

In order to identify and remove content that violates our Community Standards, we use a combination of artificial intelligence, human review, and user reports.

1. Use of Automated Tools:

Every day, we remove millions of violating pieces of content and accounts on Facebook. In most cases, this happens automatically, with technology to detect, restrict, and remove content and accounts that may go against our Community Standards and Advertising Standards. In other cases, our technology selects content for human review. Our review teams review a blend of user reports and content surfaced by our technology. Our technology also supports the review teams by prioritizing the most critical content to be reviewed, based on severity, virality, and likelihood of a violation. Our review systems use technology to prioritise high-severity content with the potential for offline harm (e.g., posts related to terrorism and suicide) and viral content that is spreading quickly and has the potential to reach a large audience, in order to prevent as much harm as possible. Our technology is trained to identify violations of our Community Standards and Advertising Standards.

There are three primary forms of technology used to detect possible policy violations. First, we employ rate limits (speed limits) on how rapidly accounts can perform multiple actions on our platforms, including making posts, to prevent the usage of bots. Next, we have matching technology that identifies identical or near-identical copies of accounts, URLs, text, images, audio, and videos that we have previously identified as violating our policies. When we match

the content exactly or we determine it is near-identical, we will typically remove the content. Finally, with appropriate oversight, we also use artificial intelligence (AI) to augment and scale our human review capacity: like with the matching technology, when confident enough that a post violates one of our Community Standards, the artificial intelligence will typically remove the content or demote it. We also use artificial intelligence to select the content for human review on the basis of severity, virality, and likelihood of a violation. As with matching technology, artificial intelligence operates on URLs, text, images, audio, and videos. Unlike technologies that can only match violations they've seen before, artificial intelligence has the potential to identify certain violations it has never seen before.

In the context of advertisements, when advertisers place an order, each ad is reviewed against our policies. Our Advertising Standards provide policy detail and guidance on the types of ad content we allow and the types of ad content we prohibit. Our Advertising Standards also provide guidance on advertiser behavior that may result in advertising restrictions being placed on a business account or its assets (i.e., an ad account, Page, or user account). Our ad review system relies primarily on automated tools to check ads and business assets against our policies.

2. Persons in Charge of Content Moderation:

Human reviewers are provided with various tools and resources when undertaking content review. For example, human reviewers receive in-depth training and often specialize in certain policy areas and regions. Human reviewers work on cases where it's essential to have people review – for example in critical situations where our technology might be unsure whether the content is violating. Reviewers may also be able to use a highlighting tool for dangerous organizations based on the region where the content is reviewed and tool-tips that explain the definitions of certain words and how they should inform decisions.

3. User reports:

Logged-in Facebook users can report Community Standards violations in a variety of ways, including through a report link available on each piece of content. People who are logged-out can also report Community Standards violations using this [form](#). When content is reported through these tools, we review the reported content (automatically or manually) to determine whether it violates our Community Standards, and we take appropriate action if it does.

If a reported piece of content is determined to violate our Community Standards, it will be promptly removed and made unavailable worldwide.

In addition to the above-mentioned methods, in 2019, we set up an independent [Oversight Board](#), bringing together experts from around the world, whose mission is to inform and advise on some of the most complex moderation decisions.

Our Advertising Standards provide policy detail and guidance on the types of ad content we allow and prohibit. Our Advertising Standards also provide guidance on advertising behavior that may result in restrictions being placed on a business account or its assets (i.e., an ad account, Page, or user account). The full set of advertising policies can be found [here](#).

B. We allow individuals and their authorized representatives in the European Union to report content that they believe violates local laws. This includes [dedicated channels](#) to report content if users believe content violates their personal legal rights or applicable local laws, [Defamation Reporting Forms](#), and [Intellectual Property Reporting Forms](#) to report copyright and trademark violations. When something on Facebook is reported to us as violating local law, but it does not go against our Community Standards, we may restrict access to the content in the jurisdiction where it is alleged to be unlawful.

C. In the event that we receive requests from judicial or administrative authorities to act against specific items of alleged illegal content on our platform, first, we review the reported content in line with our Community Standards and other relevant policies, for example our Advertising Standards. If we determine that the content goes against our policies, we remove it. If content does not go against our policies, in line with our commitments as a member of the Global Network Initiative and our Corporate Human Rights Policy, we conduct a review to confirm whether the order is valid, as well as perform human rights due diligence. We may then restrict access to the content in the jurisdiction where it is alleged to be unlawful.

2. IDENTIFYING, REMOVING, ACCESS RESTRICTION OR PREVENTING THE REAPPEARANCE OF TERRORIST CONTENT

Article 7(3)(a) information about the hosting service provider's measures in relation to the identification and removal of or disabling of access to terrorist content;

Article 7(3)(b) information about the hosting service provider's measures to address the reappearance online of material which has previously been removed or to which access has been disabled because it was considered to be terrorist content, in particular where automated tools have been used;

There is no place on Facebook for terrorism. We remove terrorists and posts that support terrorism whenever we become aware of them. Meta has one of the most comprehensive policies in the industry targeting terrorist organizations, as well as hate groups and large-scale criminal organizations, like cartels.

We [define](#) terrorist organizations and individuals as: *“any non-state actor that:*

- *Engages in, advocates or lends substantial support to purposive and planned acts of violence,*
- *Which causes or attempts to cause death, injury or serious harm to civilians, or any other person not taking direct part in the hostilities in a situation of armed conflict, and/or significant damage to property linked to death, serious injury or serious harm to civilians*
- *With the intent to coerce, intimidate and/or influence a civilian population, government or international organization*
- *In order to achieve a political, religious or ideological aim.”*

Our definition is agnostic to the ideology or political goals of a group, which means it includes everything from religious extremists and violent separatists to white supremacists and militant environmental groups. It’s about whether they use violence to pursue those goals. This definition and policy are applied globally.

Under the Community Standards, we have a set of policies that prohibit terrorist content on Facebook, including our policies on Dangerous Organizations and Individuals, Violence and Incitement, and Coordinating Harm and Promoting Crime. Taken together, these policies are congruent with the Regulation’s definition of “terrorist content.”

- Our [Dangerous Organizations and Individuals](#) policy prohibits any organization or individual that proclaims a violent mission or is engaged in violence from having a presence on Facebook. Specifically, we do not allow on our platform:
 - Terrorist organizations and individuals
 - Hate entities
 - Criminal organizations
 - Multiple-victim violence and serial murderers (including attempted multiple victim violence)
 - Hateful ideologies
 - Violent non-state actors
 - Militarized social movements
 - Violence-inducing conspiracy networks
 - Hate banned entities

Under our Dangerous Individual and Organizations policy, we designate and ban individuals and organizations tied to terrorism, organized hate and large-scale criminal activity, attempted multiple-victim violence, multiple victim violence, serial murders, and violent events. We also remove all praise, substantive support, and representation of these individuals and organizations when we become aware of it. We also designate and ban militarized social movements, violent non-state actors, and violence-inducing conspiracy networks like QAnon,

and we do not allow them to have a presence on our platform. Accounts, pages, or groups representing these organizations or individuals are removed when we're made aware of them.

- Our [Violence and Incitement](#) policy prohibits content that incites or facilitates violence and constitutes a credible threat to public or personal safety. We remove content, disable accounts, and work with law enforcement when we believe there is a genuine risk of physical harm or direct threats to public safety.
- Our [Coordinating Harm and Promoting Crime](#) policy prohibits users from facilitating, organizing, promoting, or admitting to certain criminal or harmful activities targeted at people, businesses, property, or animals. We allow people to debate and advocate for the legality of criminal and harmful activities, as well as draw attention to harmful or criminal activity that they may witness or experience as long as they do not advocate for or coordinate harm.

We detect dangerous organizations and terrorist content via a playbook and a series of automated techniques, which were initially developed to detect content related to widely known terrorist organizations like ISIS, al Qaeda, and their affiliates. We have since expanded our playbook and techniques.

Measures to identify and remove or disable access to terrorist content, individuals, or organizations include:

- We continue to designate dangerous organizations and individuals based on their behavior both online and offline – most significantly, their ties to violence. As part of the designation process, we identify signals that indicate a banned organization has a presence, and then we use technology to “fan out” and proactively investigate associated accounts, Pages, and Groups, before removing this “cluster” all at once.
- We use media-matching technology to detect when someone is trying to upload content that is identical or near-identical to photos, videos, text, and even audio that we have already removed for violating our policies around terrorism. This means that if we previously removed a propaganda video from ISIS, we can work to prevent other accounts from uploading the same video to our site. In many cases, this means that terrorist content intended for upload to Facebook or Instagram simply never reaches the platform.
- We also use the hash (unique digital fingerprints) sharing database of known violent terrorism imagery or propaganda, provided by the Global Internet Forum to Counter Terrorism (GIFCT), to help quickly detect and remove content involving a hash. The database houses hundreds of thousands of unique hashes. Again, in many cases, this means that terrorist content intended for upload never reaches our platforms.
- We use machine learning classifiers to help us understand whether content uploaded to our platforms is advocating for terrorism, based on previous content we have removed for praising or supporting terrorist organizations.

- We've increased our capability to rapidly respond to live streams that may breach our Community Standards, including by reviewing all live streams in an area that may involve footage of an attack and increasing our 24/7 capacity to respond to live stream reports. As well as our efforts around rapid response, we've also introduced [additional restrictions](#) on who can use Facebook Live. We apply a 'one strike' policy to Live in connection with a broader range of offenses. Anyone who violates our most strict policies are restricted from using Live for set periods of time – for example 30 days – starting on their first offense. For instance, someone who shares a link to a statement from a terrorist group with no context is immediately blocked from using Live for a set period of time.
- We have a team of specialists who exclusively or primarily focus on countering dangerous organizations as their core responsibility. This group includes former academics who are experts on counterterrorism, former prosecutors and law enforcement agents, investigators and analysts, and engineers. This team continues to study trends in terrorism, organized hate, and other dangerous organizations and works with partners to better understand these organizations as they evolve.
- To help other platforms that may not have the resources and the technology, we have developed and [made available](#) a free open source software tool called Hasher-Matcher-Actioner (HMA) that identifies copies of images or videos and takes action against them *en masse*. HMA builds on Meta's previous open source image and video matching software, and it can be used for any type of violating content. We hope the tool will be adopted by a range of companies to help them stop the spread of terrorist content on their platforms, and that it will be especially useful for smaller companies who don't have the same resources as bigger companies.

Measures to address the reappearance of terrorist content, individuals or organizations include:

- We take a cross-platform approach and look for banned organizations across Facebook and Instagram.
- Before removing banned terrorist groups, we investigate associated accounts, Pages, and Groups, and we remove the entire network.
- Once we remove their presence, we work to identify and disrupt attempts by the group to come back on our platform, using a tactic called a strategic network disruption. We do this by identifying signals that indicate a banned organization has a new presence and then proactively investigating associated accounts, Pages, and Groups, before removing them all at once.

Metrics demonstrating the efficacy of our efforts to reduce the prevalence of terrorist content on Facebook can be found in our quarterly [Community Standards Enforcement Report](#).

Meta recognizes its potential to impact human rights and has taken steps to demonstrate our commitment to human rights principles. We joined the Global Network Initiative, and made

commitments in our Corporate Human Rights Policy to respect human rights as set out in the United Nations Guiding Principles on Business and Human Rights (UNGP) and other international human rights standards. Human rights also guide our work developing responsible innovation practices, including when building, testing, and deploying products and services enabled by artificial intelligence (AI). We implement our commitment to human rights using approaches set out in the United Nations Guiding Principles on Business and Human Rights (UNGPs). These approaches include: (1) applying human rights policies; (2) conducting human rights due diligence and disclosure; (3) providing access to remedy; (4) maintaining oversight, governance, and accountability; and (5) protecting human rights defenders.

3. REMOVALS, ACCESS RESTRICTION AND APPEALS

Article 7(3)(c) the number of items of terrorist content removed or to which access has been disabled following removal orders or specific measures, and the number of removal orders where the content has not been removed or access to which has not been disabled pursuant to the first subparagraph of Article 3(7) and the first subparagraph of Article 3(8), together with the grounds therefor;

Article 7(3)(d) the number and the outcome of complaints handled by the hosting service provider in accordance with Article 10;

Article 7(3)(g) the number of cases in which the hosting service provider reinstated content or access thereto following a complainant by the content provider.

The data provided below reflects the volume of requests for content removal received via this channel, as well as proactive measures that Meta takes in line with its own policies, and user appeals that arise in respect of those voluntary, proactive measures.

Removals and Access Restriction

We have set up a “Regulation on Dissemination of Terrorist Content Online (EU) Removal Order Reporting Form” for Competent Authorities to submit removal orders and a process to expeditiously deal with these orders. When we receive a report, we first review it against the Facebook Community Standards. If we determine that the content goes against our policies, we remove it. If the content is not violating the Community Standards, we then conduct a review of the order and take further action as appropriate, which may include restricting access to the content in the jurisdiction where it is alleged to be unlawful.

For the period of 1 January 2023 to 31 December 2023, we received 143 requests through the dedicated channel for TCO removal orders on Facebook. However, only 15 of them were in fact valid orders from Competent Authorities designated by Member States.

Among the valid orders, 10 of them led to content removal or access restriction in the EU.

As noted in Section 1 and 2 above, we also proactively detect and take action on content that violates our Community Standards and other policies. For the period of 1 January 2023 to 31 December 2023, we removed 6.1 million pieces of content on Facebook in the EU that violated our policies on Dangerous Organizations and Individuals, Violence and Incitement, and Coordinating Harm and Promoting Crime. The number of content removals, received appeals, and prevalence rate for various categories are detailed in our [Community Standards Enforcement Report](#).

Complaints Handled (Appeals)

Where content has been removed or access thereto disabled in accordance with our policies, users have the option to report that they [disagree with our decision](#). We have a comprehensive system for processing appeals, and decisions will be overturned where a review concludes that is the appropriate outcome.

For the period of 1 January 2023 to 31 December 2023, we received 857.2 thousand appeals from users in the EU related to content on Facebook that we removed due to violations of our policies on Dangerous Organizations and Individuals, Violence and Incitement, and Coordinating Harm and Promoting Crime. We restored 156 thousand pieces of content following those appeals.

No complaints were handled by Meta in accordance with Article 10 under EU TCO during the reporting period. There are also no cases in which content or access to content has been reinstated further to a complaint under Article 10 of EU TCO.

4. ADMINISTRATIVE OR JUDICIAL REVIEW PROCEEDINGS

Article 7(3)(e) the number and the outcome of administrative or judicial review proceedings brought by the hosting service provider;

Article 7(3)(f) the number of cases in which the hosting service provider was required to reinstate content or access thereto as a result of administrative or judicial review proceedings.

For the period of 1 January 2023 to 31 December 2023, no administrative or judicial review proceedings were issued. On this basis, we have no further information to include with regard to Article 7(3)(e) or (f).

Note:

1. The metrics in this report are our best estimates of content we took action on based on the creator of the content and predicted country locations for those users.
2. The metrics in this report are based on definitions and caveats as disclosed in [Content Actioned](#).
3. Given that such violations are also highly adversarial, EU-level data, which is an aggregate of country-level data derived from the 27 EU Member States, including their outermost regions, may be less reliable. For example, bad actors may often try to avoid detection by our systems by masking the country they are coming from. While our enforcement systems are global and will try to account for such behavior, this makes it very difficult to attribute and report the accounts or content by producer country (where the person who posted content was located).
4. This report shares metrics in which we have considerable confidence in their accuracy on Facebook.