



Türkiye Transparency Report for Content Removal Requests under Law No. 5651

July 31, 2024

The following report is published in accordance with Additional Article 4 (4) of Türkiye's Law No. 5651 on Regulation of Broadcasts via Internet and Prevention of Crimes Committed through Such Broadcasts ('**Law No. 5651**') and Procedures and Principles Regarding Social Network Providers. This report contains information for a period of six months, from January 1, 2024 to June 30, 2024.

1. GENERAL OBSERVATIONS¹

Overview of Facebook's content moderation policies and reporting forms.

Meta Platforms, Inc., which provides the social network service "Facebook" for users in Türkiye, recognizes how important it is for Facebook to be a platform where people feel empowered to communicate and express themselves, and we take our role to maintain a safe and trusted platform seriously.

A. We maintain a robust set of policies, including our [Community Standards](#) and [Advertising Standards](#) that define what is and is not allowed on Facebook. These standards apply to content worldwide and are integral to protecting both expression and personal safety on Facebook.

Facebook is a global community committed to equal treatment of all users. The [Facebook Community Standards](#) are applied consistently to everyone on the platform, irrespective of their

¹ This report contains various links to external documents including:

- Meta's Transparency Center. Pages in the Transparency Center are viewable in Turkish by toggling to the Turkish Language option that appears in the top right corner of the [Transparency Center home page](#).
- Facebook Help Center. Pages in the Transparency Center are viewable in Turkish by toggling to the Turkish Language option that appears in the top right corner of the [Help Center](#).
- Meta's [Newsroom](#).

location, and to all types of content. These policies are based on feedback from people and the advice of experts in fields such as technology, public safety, and human rights.

Our reviewers enforce these standards using comprehensive guidelines, which ensure that decisions are as consistent as possible and that users are treated equally and impartially.

Our [Community Standards](#) prohibit a wide range of objectionable or harmful content, including content that:

- Violates someone's privacy, personal information, and/or image privacy rights
- Is considered bullying or harassment
- Incites or facilitates violence and credible threats to public or personal safety
- Is considered hate speech
- Is considered graphic or violent imagery
- Is considered spam

Meta's Content Policy team, which sits in more than a dozen locations around the world, is responsible for developing our Community Standards. The team includes subject matter experts on issues such as hate speech, child safety, and terrorism, as well as people with experience in criminal prosecution, rape crisis counselling, academics, human rights, law, and education. Many have also worked on issues of voice and safety long before coming to Meta.

This team regularly seeks input from outside experts and organizations to help balance the different perspectives that exist on free expression and safety, and to better understand the impacts of our policies on different communities globally.

Our [Advertising Standards](#) provide policy detail and guidance on the types of ad content we allow and the types of ad content we prohibit. Our Advertising Standards also provide guidance on advertiser behavior that may result in advertising restrictions being placed on a business account or its assets (an ad account, Page, or user account).

In order to identify and remove content that violates our Community Standards, we use a combination of machine learning, human review, and user reports.

1. Use of Automated Tools:

Every day, we remove millions of violating pieces of content and accounts on Facebook. In most cases, this happens automatically, with technology to detect, restrict, and remove content and accounts that may go against our Community Standards and Advertising Standards. In other



cases, our technology selects content for human review. Our review teams review a blend of user reports and content surfaced by our technology. Our technology also supports the review teams by prioritizing the most critical content to be reviewed, based on severity, virality, and likelihood of a violation. Our review systems use technology to prioritize high-severity content with the potential for offline harm (e.g., posts related to terrorism and suicide) and viral content that is spreading quickly and has the potential to reach a large audience, in order to prevent as much harm as possible. Our technology is trained to identify violations of our Community Standards and Advertising Standards.

There are three primary forms of technology used to detect possible policy violations. First, we employ rate limits (speed limits) on how rapidly accounts can perform multiple actions on our platforms, including making posts, to prevent the usage of bots. Next, we have matching technology that identifies identical or near identical copies of accounts, URLs, text, images, audio, and videos that we have previously identified as violating our policies. When we match the content exactly or we determine it is near identical, we will typically remove the content. Finally, we also use machine learning (ML) to augment and scale our human review capacity with appropriate oversight: like with the matching technology, when confident enough that a post violates one of our Community Standards, the machine learning will typically remove the content or demote it. We also use machine learning to select the content for human review on the basis of severity, virality, and likelihood of a violation. As with matching technology, machine learning operates on URLs, text, images, audio, and videos. Unlike technologies that can only match violations they've seen before, machine learning has the potential to identify certain violations it has never seen before.

In the context of advertisements, when advertisers place an order, each ad is reviewed against our policies. Our ad review system relies primarily on automated tools to check ads and business assets against our policies.

2. Persons in Charge of Content Moderation:

Human reviewers are provided with various tools and resources when undertaking content review. For example, human reviewers receive in-depth training and often specialize in certain policy areas and regions. Please see Section 8 below on human resources dedicated to content moderation, which includes details on the measures taken to provide training and assistance to persons in charge of content moderation. For example, reviewers may be able to use a highlighting tool for slurs and dangerous organizations based on the region where the content is reviewed and tooltips that explain the definitions of certain words and how they should be used to inform decisions.

3. User Reports and Removal Requests:

A. Logged-in Facebook users can report Community Standards violations in a variety of ways, including through a report link available on each piece of content. When content is reported through these tools, we review the reported content (automatically or manually) to determine whether it violates our Community Standards and take appropriate action if it does.

If a reported piece of content is determined to violate our Community Standards, it will be promptly removed and made unavailable worldwide.

In addition to the above-mentioned methods, we have set up an independent [Oversight Board](#), bringing together experts from around the world, whose mission is to inform and advise on some of the most complex moderation decisions.

Our Advertising Standards provide policy detail and guidance on the types of ad content we allow and prohibit. Our Advertising Standards also provide guidance on advertising behavior that may result in restrictions being placed on a business account or its assets (i.e., an ad account, Page, or user account). The full set of advertising policies can be found [here](#).

B. Under our policies, rights holders can submit intellectual property reports via [Intellectual Property Reporting Forms](#), identifying content on Facebook that they believe infringes their rights. If a report is complete and valid, we will act promptly to remove the reported content.

C. Pursuant to the Law No. 5651, everyone in Türkiye (users and non-users) can report content on Facebook that they believe violates Articles 9 and 9/A of the Law No. 5651 using a specific [reporting form](#). This form also informs users of how to report other issues that they encounter on Facebook, including hacked accounts, impersonation, privacy violations, and fake accounts. Users in Türkiye can access this form through a dedicated [Facebook Help Center Page](#).

D. In the event that we receive requests from judicial or administrative authorities to act against specific items of alleged illegal content on our platform, first, we review the reported content in line with our Community Standards and other relevant policies, for example our Advertising Standards. If we determine that the content goes against our policies, we remove it. If content does not go against our policies, in line with our commitments as a member of the Global Network Initiative, we conduct a careful legal review to confirm whether the report is valid, as well as a human rights due diligence. We may then restrict access to the content in the jurisdiction where it is alleged to be unlawful.

E. In cases where we assess that reports are not legally valid, are overly broad, or are inconsistent with international human rights standards, we may request clarification or take no action. There also may be other reasons why access to content has not been removed or restricted, such as no URL was provided, or the provided URL is not leading to allegedly violating content, or the content has already been removed.

More information on our approach to content moderation can be found in the [Meta Transparency Center](#).

2. REPORTING MECHANISMS/CRITERIA

Description of the mechanisms for submitting requests under the Law No. 5651.

A. INFORMATION REGARDING APPLICATIONS FROM INDIVIDUALS

People in Türkiye can submit applications under the Law No. 5651 using Facebook’s dedicated [reporting form](#). In order to evaluate the application, people are prompted to provide the following information in our dedicated application form:

- Complete contact information, including whether the applicant submits the application on behalf of themselves/their organization or is an attorney or other representative authorized to report on behalf of a client
- Links to the specific content on Facebook that is the subject of the report
- Section(s) of the Law No. 5651 under which the application is made
- Why these statements or images are alleged to violate Articles 9 and/or 9/A of Law No. 5651
- The applicant’s full name

Everyone in Türkiye can access this [form](#) from our dedicated [Help Center Page](#) and submit an application by clicking “Send” once all required information is filled out in the Form. In addition, the Help Center Page is searchable via the Help Center search bar.

Once the applicant has completed the form described above and clicks “send,” their application automatically reaches us through an internal review tool – no additional action is required by the reporting party.

B. INFORMATION REGARDING APPLICATIONS FROM JUDICIAL AND ADMINISTRATIVE AUTHORITIES

Judicial and administrative authorities from Türkiye can submit applications under Law No. 5651 through dedicated communication channels. Upon receipt of an application by judicial and administrative authorities from Türkiye, all the information from that application such as validity and description/identification of the content in question is reviewed.

3. INFORMATION RELATED TO CONTENT REMOVAL REQUESTS

Statistical and categorical information about the reporting channel and summary of our actions.

A. APPLICATIONS RECEIVED FROM INDIVIDUALS

	Volume
Total number of applications	203
Total number of duplicative applications	59
Total number of reported entities	376
Total number of reported large object entities ²	340
Total number of other objects reported ³	36
Total number of removed entities	17

- Between January 1, 2024 and June 30, 2024, Meta Platforms, Inc. has received a total of 203 applications through the dedicated Facebook reporting form.
- 59 of these applications were duplicative of prior applications, i.e., the same entity was reported multiple times.
- 340 of the 376 total number of entities reported were “large objects,” i.e., entire Facebook profiles, Pages, or groups, rather than specific pieces of content such as photos, status, or posts.
- After making an assessment, we concluded that 17 of these 376 reported entities were violating and were removed from our platform for violating our policies. The top reasons for removing these reported entities included violations against our private impersonation, privacy, fraud and bullying and harassment policies.

² Large object entities include Facebook profiles, Pages and groups.

³ Includes all reported small pieces of content from Facebook and any type of Instagram entities reported through the Facebook dedicated reporting form.

B. REQUESTS RECEIVED FROM JUDICIAL AND ADMINISTRATIVE AUTHORITIES

	Volume
Total number of requests	2,953
Total number of ICTA⁴ requests⁵	1849
Total number of Consumer Policy requests	24
Total number of court orders	463
Total number of reported entities	2,252
Total number of entities removed for violating policies	792
Total number of entities restricted in Türkiye	920

- Between January 1, 2024 and June 30, 2024, Meta Platforms, Inc. has received a total of 2953 removal requests from judicial and administrative authorities.
 - 1849 of these requests were received from the Information and Communication Technologies Authority (Article 8/8A of Law No. 5651), in some instances containing both Facebook and Instagram content together;
 - 24 of these requests were received from the Pharmaceuticals and Medical Devices Administration of Türkiye and the Board of Advertisement of Türkiye through the Consumer Policy Channel;
 - 463 of these requests were court orders received from the Internet Access Providers Union.
- These removal requests contained a total of 2252 reported entities.
- After making a careful assessment,
 - 792 of the reported entities were removed for violating our policies;
 - 920 of the reported entities were restricted in Türkiye on the basis of local law.

⁴ The Information and Communication Technologies Authority.

⁵ The metrics presented herein encompass the aggregate count of both Facebook and Instagram reports. It is noteworthy that the reports submitted to us in some instances contain both Facebook and Instagram content together.



4. USER CONTROLS

Summary of policies, tools, products, and resources available to users to keep them safe on Facebook

We believe that users should be empowered to customize their online experience and be given [tools and resources](#) to help keep themselves safe on the platform.

- **Ranking Transparency & Controls:** We want people to better understand why they see certain posts and ads in their Feed and to be able to decide if they want to continue seeing those posts or ads. Since 2014, we have introduced tools to explain how users' past interactions impact the ranking of posts in their Feed. A full list of ranking tools, products, and resources can be found [here](#).
 - [Why am I seeing this post?](#) helps people understand and more easily control what they see from friends, Pages, and Groups in their Feed. It explains how people's activities impact the ranking of content in their Feeds (e.g., if the post is from a friend, Page, or Group the user follows), as well as what other information influences the order of posts (e.g., how often the user interacts with a specific type of post like videos, photos, links). This tool also provides shortcuts to controls, such as See First, Unfollow, Feed Preferences, and Privacy Shortcuts, to help users personalize their Feed.
 - [Why am I seeing this ad?](#) allows people to see how factors like basic demographic details, interests and website visits contribute to the ads that are shown in their Feeds. There are also additional details about when information on an advertiser's audience list matches a person's profile.
- **Feed Ranking Controls:** Users can control what they see in their Feed through the following tools.
 - [Feed Preferences](#) provides a suite of tools that allow people to manage what they see in their Facebook Feed, including the ability to unfollow people, snooze a particular account, or prioritize Favorites.
 - [Favorites Feed](#) allows users to control and prioritize posts from friends and Pages they care about most. By selecting up to 30 friends and Pages to include in Favorites, their posts will appear higher in ranked Feed and can also be viewed in a separate feed populated exclusively with posts from a user's



“Favorites.”

- [Most Recent Feed](#) allows users to see content sorted in their Feed by chronological order with the newest post first.
- [Feeds Tab](#) allows users to alternate between different Feed experiences: the algorithmically ranked Top Posts Feed, the chronological Most Recent Feed, or the Favorites Feed.
- [Reduce Control](#) allows users to control whether or how strongly certain types of content are demoted on Feed. Demotions can be left as-is, made stronger, or turned off. Turning off demotions is age-restricted and unavailable to Facebook users whose age setting is below their country’s age of majority.

5. ALGORITHM MECHANISMS/CRITERIA

Explanation on how the algorithm works in Feed.

The posts that users see in Feed are ranked based on what we believe will be most valuable to users. Feed is largely comprised of posts from the friends, groups, and Pages users have chosen to connect with. Because most people have more content in their Feed than they could possibly browse in one session, we use an algorithm to determine the order of all the posts users could see. To prioritize the most meaningful posts at the top of your Feed, the algorithm works in four steps:

- **Inventory:** The first item the algorithm considers is the total set of posts a user could see when they open Facebook. This includes all the posts shared by the people the user connected to as 'friends,' the Pages that the user follows and the groups that they have joined, interspersed with ads, and recommended content that we think will be relevant to users based their Facebook activity.
- **Signals:** Then, for each of these posts, the algorithm considers multiple factors such as who posted it, how the user has previously interacted with that person, whether it's a photo, a video, or a link, and how popular the post is, based on things such as how many of the users friends liked it and Pages that re-shared it. All these factors are called signals.
- **Prediction:** From there, the algorithm uses these signals to make a series of personalized predictions about each post based on how likely it is to be relevant to a user: for example, whether it's from your friends or family, how likely you might be to comment on it, how likely it is to foster a [meaningful interaction](#), how likely it might be for a user to find it on their own



or if it contains a quality indicator (if a piece of news is [original content](#), the algorithm assigns it a higher personalized relevance score, and it will often appear closer to the top of a user's Feed). We also run a number of surveys asking people whether a post was "[worth your time](#)," and based on those survey responses, we predict how likely people are to find a post worthwhile. Posts that are predicted to be more worthwhile are shown higher up in Feed.

- **Relevance:** Lastly, the algorithm calculates a relevance score for each post in a user's inventory based on these signals and predictions. Posts with higher scores are more likely to be interesting to a user, so they'll be placed closer to the top of the user's Feed, and posts with lower scores will be closer to the bottom.

Conversely, there are some kinds of posts that people have told us they don't want to see or are broadly understood to be harmful, so we strive to remove content from Facebook altogether when it poses a real risk of harm, such as graphic violence, hate speech, or fake COVID cures. Therefore, Facebook has [Community Standards](#) that prohibit harmful content, and this is why we invest heavily in developing ways of identifying and acting on it quickly. We also use Feed ranking to [reduce the distribution](#) of posts that may contain content that people find objectionable but don't necessarily meet the bar of removal under our policies. If a post is likely to contain misinformation, a sensationalized health claim or clickbait, for example, it will receive a lower value score and appear lower in Feed as a result. We monitor for new types of problematic content such as these and develop new algorithmic levers to detect and enforce against them. However, we are confronted by the challenges that come with a platform used by an ever-growing population. As such, our teams work hard every day to adapt and address in a thoughtful and deliberate way any gaps that may arise in a thoughtful and deliberate way.

We have also recently announced the publication of [22 system cards](#), providing an unprecedented level of insight into how our AI systems rank content on Facebook and Instagram. These cards provide information about how our AI systems rank content for Feed, Reels, Stories, and other surfaces; some of the predictions each system makes to determine what content might be most relevant to people; and the options available to help customize an experience on Facebook. These build on our long-standing "[Why Am I Seeing This](#)" feature, which allows people to see details directly in our applications about why our systems predicted that specific content would be relevant to them, and the types of activity and inputs that may have led to that prediction. Furthermore, in June 2023, we've also updated [Our Approach to Feed Ranking](#) to help people maximize their experience and help creators understand how their content might be surfaced. In parallel, we have created a dedicated space on the topic of [demotion](#) across our platforms in our Transparency Center.

6. ADVERTISING TRANSPARENCY

We have provided information about our advertising policies in Section 1 (General Observations), which sets out the types of ads content we allow or prohibit. In addition to these policies, we maintain an [ad library](#) that enhances advertising transparency by providing a comprehensive searchable database of all ads currently running on Facebook, along with dates the ad ran, the parameters used for targeting (e.g., age, gender, location), who was served the ad, and more.

Additionally, as set out in Section 4 (Information on User Controls), when users see an ad, they can select [Why am I seeing this ad?](#) to view details on why the user is seeing a specific ad in their feed.

7. AUTOMATED MEANS FOR CONTENT MODERATION

Any use made of automated means for the purpose of content moderation.

A. Use of automated means for the purpose of own initiative and other content moderation, and purpose of those tools

As described in Section 1 (General Observations), we use technology to help us proactively detect content on our services that might be harmful and violate our Community Standards, so we can remove it faster. We described in Section 5 our primary content moderation automated tools.

These technologies run on accounts, posts, comments, photos, and other pieces of content uploaded to Facebook. They determine how probable or likely it is that this content violates a certain policy, based on those signals or patterns, and if the content should be automatically removed.

B. Indicators of accuracy, error rates, safeguards

Our technology learns and improves from each human decision. Over time – after learning from thousands of human decisions – the technology gets better. When reviewing violating content, review teams manually label the policy guiding their decision, which means that they mark or “label” the relevant policy that the content, account, or behavior violates. This labeling of data helps us improve the quality of our algorithms that proactively detect and remove harmful content, accounts, and behavior.

To ensure and improve the quality, i.e., how accurate the technologies mentioned above are in enforcing Community Standards and other policy violations, there are ongoing quality evaluation processes in place. Meta uses overlapping techniques and systems for maintaining a high overall accuracy for our automated content moderation.

Prior to fully launching any new rate limit (speed limit), matching technology or machine learning (ML), we use the technology to only log how the technology would have behaved instead of immediately acting. We then use human reviewers to assess the accuracy against current content, behavior, or accounts, rather than just historical ones, as we did during the technology's training. After launching rate limits, matching technologies, or machine learning, we monitor the volumes of actions and appeals by the user who posted the content as well as the rate at which appeals are granted. If any of the metrics we monitor are abnormal, our engineering teams may investigate.

For each primary form of automation technology, the investigation of abnormal metrics can vary. With rate limits, engineers typically reevaluate if the limit is preventing bot behavior. For our matching technologies, if an entry in our list of previously identified instances of policy violations has abnormal signals, we will re-review the entry to confirm it continues to go against our policies. Similarly, if one of our machine learning classifiers has abnormal signals, we will either send a sample of the machine learning classifier's recent results to human labeling to confirm the accuracy rate or deprecate the classifier if abnormal signals indicate a clear breakage.

In addition, many of our machine learning classifiers are automatically reassessed for accuracy after each human review. This classifier reassessment is an example of the general feedback loop between human review and technology. The content labeling decisions taken by human reviewers are used to train and refine our technology. As a part of this process, the review teams manually label the policy guiding their decision, i.e., they mark the policy that the content, account, or behavior violates. This helps to improve the quality of our machine learning algorithms and our lists of known policy-violating content used by our matching technology. To maintain quality control in all of these decisions, we regularly audit random samples of decisions taken by the algorithm and our content reviewers and measure them against our expectations for policy enforcement. In the context of automation relating to language, some automation is developed to support specific languages whilst others are language agnostic.

8. PERSONS IN CHARGE OF CONTENT MODERATION

The human resources that Meta dedicates to content moderation in respect of the service offered in Türkiye, their qualifications and linguistic expertise and training and support given to such staff.

As described in Section 1 (General Observations), Meta uses human resources dedicated to content moderation, including for compliance with the obligations set out in Articles 8 and 8/A and Article 9 & 9/A of the Law No. 5651.

A. Qualifications And Training

Human reviewers come from different backgrounds, reflect our diverse community, and include experts in enforcement in policy areas such as child safety, hate speech, and counterterrorism. Human reviewers undergo extensive training when they join and are regularly trained and tested beyond this initial training, with specific examples, such as on how to uphold the Community Standards and take the correct action on a report. We also do our own proactive audits, where we conduct re-reviews that help us figure out if we are getting it right.

Human reviewers who review content alleged to be illegal receive distinct training based on the nature of their respective work. Every member receives several weeks of training focused heavily on operational proficiency and in preparation for processing such content. For example, the reviewers who review content for defamation receive training specifically on assessing defamation, including background material on the law.

B. Support

We recognize that reviewing content can be challenging work. Keeping people safe online sometimes means review teams have to look at content that may be objectionable or graphic. We respect the difficulty of this work and work with industry leading vendors to ensure reviewers have access to the resources they need to do their job and support their health. There is a robust and diverse program to support human reviewers. Our vendor contracts mandate high quality support in a variety of areas, including pay, benefits, work environment, and wellbeing and psychological support. The assistance model depends on what type of content reviewers work on. Such assistance, for example, can take the form of psychological support, including individual and group sessions, and a 24/7 independent support program which includes a range of offerings that include clinical services.

C. Volume of human resources dedicated to content moderation and linguistic capabilities

The team working on safety and security is made up of around 40,000 people. About 15,000 of those are content reviewers, and they include a mixture of full-time employees, contractors, and outsourced support. We partner with companies to help with content review, which allows us to scale globally with coverage across time zones, languages, and markets.

This content review team is global and reviews content 24/7 in over 80 languages, including Turkish. The team includes reviewers with language expertise to enforce our policies in cases where certain words or content require additional contextual understanding. Not all content requires language expertise; for example, nudity and sexual activity is language agnostic. For

this type of content, Meta has a global pool of content moderators who review these types of reports. Currently, there are around 2,500 language agnostic reviewers globally. These reviewers also review Turkish content that falls under this language category.