



Türkiye Transparency Report for Content Removal Requests under Law No. 5651

July 31, 2024

The following report is published in accordance with Additional Article 4 (4) of Türkiye's Law No. 5651 on Regulation of Broadcasts via Internet and Prevention of Crimes Committed through Such Broadcasts ('**Law No. 5651**') and Procedures and Principles Regarding Social Network Providers. This report contains information for a period of six months, from January 1, 2024 to June 30, 2024.

1. GENERAL OBSERVATIONS¹

Overview of Instagram's content moderation policies and reporting forms.

Meta Platforms, Inc., which provides the social network service "Instagram" for users in Türkiye, recognizes how important it is for Instagram to be a platform where people feel empowered to communicate and express themselves, and we take our role to maintain a safe and trusted platform seriously.

A. We maintain a robust set of policies, including our [Community Guidelines](#) and [Advertising Standards](#) that define what is and is not allowed on Instagram. These standards apply to content worldwide and are integral to protecting both expression and personal safety on Instagram.

¹ This report contains various links to external documents including:

- Meta's Transparency Center. Pages in the Transparency Center are viewable in Turkish by toggling to the Turkish Language option that appears in the top right corner of the [Transparency Center home page](#).
- Instagram Help Center. Pages in the Help Center are viewable in Turkish by toggling to the Turkish Language option that appears in the top right corner of the [Help Center](#).
- Meta's [Newsroom](#).

Instagram is a global community committed to equal treatment of all users. The [Instagram Community Guidelines](#) are applied consistently to everyone on the platform, irrespective of their location, and to all types of content. These policies are based on feedback from people and the advice of experts in such fields as technology, public safety, and human rights.

Our reviewers enforce these standards using comprehensive guidelines, which ensure that decisions are as consistent as possible and that users are treated equally and impartially.

Our [Community Guidelines](#) prohibit a wide range of objectionable or harmful content, including content that:

- Violates someone's privacy, personal information, and/or image privacy rights
- Is considered bullying or harassment
- Incites or facilitates violence and credible threats to public or personal safety
- Is considered hate speech
- Is considered graphic or violent imagery
- Is considered spam

Meta's Content Policy team, which sits in more than a dozen locations around the world, is responsible for developing our Community Guidelines. The team includes subject matter experts on issues such as hate speech, child safety, and terrorism, as well as people with experience in criminal prosecution, rape crisis counselling, academics, human rights, law, and education. Many have also worked on issues of voice and safety long before coming to Meta.

This team regularly seeks input from outside experts and organizations to help balance the different perspectives that exist on free expression and safety, and to better understand the impacts of our policies on different communities globally.

Our [Advertising Standards](#) provide policy detail and guidance on the types of ad content we allow and the types of ad content we prohibit. Our Advertising Standards also provide guidance on advertiser behavior that may result in advertising restrictions being placed on a business account or its assets (an ad account, Page, or user account).

In order to identify and remove content that violates our Community Guidelines, we use a combination of machine learning, human review, and user reports.

1. Use of Automated Tools:

Every day, we remove millions of violating pieces of content and accounts on Instagram. In most cases, this happens automatically, with technology to detect, restrict, and remove content and accounts that may go against our Community Guidelines and Advertising Standards. In other cases, our technology selects content for human review. Our review teams review a blend of user reports and content surfaced by our technology. Our technology also supports the review teams by prioritizing the most critical content to be reviewed, based on severity, virality, and likelihood of a violation. Our review systems use technology to prioritize high-severity content with the potential for offline harm (e.g., posts related to terrorism and suicide) and viral content that is spreading quickly and has the potential to reach a large audience, in order to prevent as much harm as possible. Our technology is trained to identify violations of our Community Guidelines, Advertising Standards.

There are three primary forms of technology used to detect possible policy violations. First, we employ rate limits (speed limits) on how rapidly accounts can perform multiple actions on our platforms, including making posts, to prevent the usage of bots. Next, we have matching technology that identifies identical or near identical copies of accounts, URLs, text, images, audio, and videos that we have previously identified as violating our policies. When we match the content exactly or we determine it is near identical, we will typically remove the content. Finally, we also use machine learning (ML) to augment and scale our human review capacity with appropriate oversight: like with the matching technology, when confident enough that a post violates one of our Community Guidelines, the machine learning will typically remove the content or demote it. We also use machine learning to select the content for human review on the basis of severity, virality, and likelihood of a violation. As with matching technology, machine learning operates on URLs, text, images, audio, and videos. Unlike technologies that can only match violations they've seen before, machine learning has the potential to identify certain violations it has never seen before.

In the context of advertisements, when advertisers place an order, each ad is reviewed against our policies. Our ad review system relies primarily on automated tools to check ads and business assets against our policies.

2. Persons in Charge of Content Moderation:

Human reviewers are provided with various tools and resources when undertaking content review. For example, human reviewers receive in-depth training and often specialize in certain policy areas and regions. Please see Section 8 below on human resources dedicated to content moderation, which includes details on the measures taken to provide training and assistance to persons in charge of content moderation. For example, reviewers may be able to use a highlighting tool for slurs and dangerous organizations based on the region where the content is reviewed and tooltips that explain the definitions of certain words and how they should be used to inform decisions.

3. User Reports and Removal Requests:

A. Logged-in Instagram users can report Community Guidelines violations in a variety of ways, including through a report link available on each piece of content. When content is reported through these tools, we review the reported content (automatically or manually) to determine whether it violates our Community Guidelines and take appropriate action if it does.

If a reported piece of content is determined to violate our Community Guidelines, it will be promptly removed and made unavailable worldwide.

In addition to the above-mentioned methods, we have set up an independent [Oversight Board](#), bringing together experts from around the world, whose mission is to inform and advise on some of the most complex moderation decisions.

Our Advertising Standards provide policy detail and guidance on the types of ad content we allow and prohibit. Our Advertising Standards also provide guidance on advertising behavior that may result in restrictions being placed on a business account or its assets (i.e., an ad account, Page, or user account). The full set of advertising policies can be found [here](#).

B. Under our policies, rights holders can submit intellectual property reports via [Copyright](#) and [Trademark](#) Reporting Forms identifying content on Instagram that they believe infringes their rights. If a report is complete and valid, we will act promptly to remove the reported content.

C. Pursuant to the Law No. 5651, everyone in Türkiye (users and non-users) can report content on Instagram that they believe violates Articles 9 and 9/A of the Law No. 5651 using a specific [reporting form](#). This form also informs users of how to report other issues that they encounter on Instagram, including hacked accounts, impersonation, privacy violations, and fake accounts. Users in Türkiye can access this form through a dedicated [Instagram Help Center Page](#).

D. In the event that we receive requests from judicial or administrative authorities to act against specific items of alleged illegal content on our platform, first, we review the reported content in line

with our Community Guidelines and other relevant policies, for example our Advertising Standards. If we determine that the content goes against our policies, we remove it. If content does not go against our policies, in line with our commitments as a member of the Global Network Initiative, we conduct a careful legal review to confirm whether the report is valid, as well as a human rights due diligence. We may then restrict access to the content in the jurisdiction where it is alleged to be unlawful.

E. In cases where we assess that reports are not legally valid, are overly broad, or are inconsistent with international human rights standards, we may request clarification or take no action. There also may be other reasons why access to content has not been removed or restricted, such as no URL was provided, or the provided URL is not leading to allegedly violating content, or the content has already been removed.

More information on our approach to content moderation can be found in the [Meta Transparency Center](#).

2. REPORTING MECHANISMS/CRITERIA

Description of the mechanisms for submitting requests under the Law No. 5651.

A. INFORMATION REGARDING APPLICATIONS FROM INDIVIDUALS

People in Türkiye submit applications under the Law No. 5651 using Instagram's dedicated [reporting form](#). In order to evaluate the application, people are prompted to provide the following information in our dedicated application form:

- Complete contact information, including whether the applicant submits the application on behalf of themselves/their organization or is an attorney or other representative authorized to report on behalf of a client;
- Links to the specific content on Instagram that is the subject of the report;
- Section(s) of the Law No. 5651 under which the application is made;
- Why these statements or images are alleged to violate Articles 9 and/or 9/A of the Law No. 5651;

- The applicant's full name.

Everyone in Türkiye can access this [form](#) from our dedicated [Help Center Page](#) and submit an application by clicking "Send" once all required information is filled out in the Form. In addition, the Help Center Page is searchable via the Help Center search bar.

Once the applicant has completed the form described above and clicks "send," their application automatically reaches us through an internal review tool – no additional action is required by the reporting party.

B. INFORMATION REGARDING APPLICATIONS FROM JUDICIAL AND ADMINISTRATIVE AUTHORITIES

Judicial and administrative authorities from Türkiye can submit applications under Law No. 5651 through dedicated communication channels. Upon receipt of an application by judicial and administrative authorities from Türkiye, all the information from that application such as validity and description/identification of the content in question is reviewed.

3. INFORMATION RELATED TO CONTENT REMOVAL REQUESTS

Statistical and categorical information about the reporting channel and summary of our actions.

A. APPLICATIONS RECEIVED FROM INDIVIDUALS

	Volume
Total number of applications	943
Total number of duplicative applications	312
Total number of reported entities	748
Total number of reported profiles	460

Total number of other entities reported²	288
Total number of removed entities	68

- Between January 1, 2024 and June 30, 2024, Meta Platforms, Inc. has received a total of 943 applications through the dedicated Instagram reporting form.
- 312 of these applications were duplicative of prior applications, i.e., the same entity was reported multiple times - we therefore have excluded these applications from the numbers below.
- 460 of the 748 total entities reported were entire Instagram profiles rather than specific pieces of content.
- After making an assessment, we concluded that 68 entities reported out of these 748 reported entities were violating and were removed from our platform for violating our policies. The top reasons for removing these pieces of content included violations of spam, scam, private impersonation policies and Advertising Standards.

B. REQUESTS RECEIVED FROM JUDICIAL AND ADMINISTRATIVE AUTHORITIES

	Volume
Total number of requests	2,953
Total number of ICTA³ requests⁴	1849
Total number of Consumer Policy requests	97
Total number of court orders	520

² Includes all reported small pieces of content from Instagram and any type of Facebook entity reported through the Instagram dedicated reporting form.

³ the Information and Communication Technologies Authority

⁴ The metrics presented herein encompass the aggregate count of both Facebook and Instagram reports. It is noteworthy that the reports submitted to us in some instances contain both Facebook and Instagram content together.

Total number of reported entities	2,580
Total number of entities removed for violating policies	504
Total number of entities restricted in Türkiye	1,941

- Between January 1, 2024 and June 30, 2024, Meta Platforms, Inc. has received a total of 2953 removal requests from judicial and administrative authorities.
 - 1849 of these requests requests were received from the Information and Communication Technologies Authority (Article 8/8A of Law No. 5651), in some instances containing both Facebook and Instagram content together;
 - 97 of these requests were received from the Pharmaceuticals and Medical Devices Administration of Türkiye and the Board of Advertisement of Türkiye through the Consumer Policy Channel;
 - 520 of these were court orders received from the Internet Access Providers Union.
- These removal requests contained a total of 2580 entities.
- After making a careful assessment,
 - 504 of the reported entities were removed for violating our policies;
 - 1941 of the reported entities were restricted in Türkiye on the basis of local law.

4. USER CONTROLS

Summary of policies, tools, products, and resources available to users to keep them safe on Instagram

We believe that users should be empowered to shape their online experience and be given [tools and resources](#) to help keep themselves safe on the platform.

- **Ranking Transparency & Controls:** We want people to better understand why they see certain posts and ads in their Feed and to be able to decide if they want to continue seeing those posts or ads. We have introduced tools to explain how user's past interactions impact the ranking of posts in their Feed. A full list of ranking tools, products, and resources can be found [here](#).
 - [Why am I seeing this post?](#) tapping (Android) or (iPhone) above a post, then tapping Why you're seeing this post allows a user to see why a certain post appears in their feed.
 - [Why am I seeing this ad?](#) Instagram wants users to see ads from businesses that are interesting and relevant to them. Users are able to see why they are seeing a specific ad by tapping the three dot menu in the top right of an ad in their feed. To adjust what ads they see, users have options to:
 - View and adjust [ad preferences](#).
 - [Hide or report an ad](#) the user does not like.
- **Feed Ranking Controls:** Users can control what they see in their Feed through the following tools.
 - [Favorites Feed](#) Users can add up to 50 accounts to their favorites allowing them to prioritize posts from those accounts. A separate feed populated exclusively with posts from a user's "Favorites" can also be accessed by tapping Instagram in the top left, then tapping Favorites.
 - [Follower Feed](#) shows posts from people a user follows on Instagram. Users can access the feed by tapping Instagram in the top left of their feed, then tap Following to find your following view.

5. ALGORITHM MECHANISM/CRITERIA

Explanation on how the algorithm works in Feed.

The posts that users see in Feed are ranked based on what we believe will be most valuable to users. Feed is largely comprised of posts from the accounts users have chosen to connect with. Because most people have more content in their Feed than they could possibly browse in one session, we use an algorithm to determine the order of all the posts users could see. With any ranking algorithm, how it works can be broken down into steps:

- **Inventory:** We start by defining the set of things we plan to rank in the first place. With Feed we consider recent posts shared by the people you follow, as well as posts from accounts you don't already follow that we think you might be interested in. We determine what you might be interested in based on a variety of factors, including what and whom you've followed, liked or engaged with recently. We personalize the experience for users to try to strike a balance between content from accounts they follow with content from accounts you don't follow but might be interested in.
- **Signals:** Next we take all the information we have about what was posted, the people who made those posts, and your preferences. We've also started considering other factors like format, so if we notice you prefer photos, we'll show you more photos. We call these "signals," and there are thousands of them. They include everything from when a post was shared to whether you're using a phone or the web to how often you like videos. The most important signals across Feed, roughly in order of importance, are:
 - Your activity. Posts you've liked, shared, saved or commented on help us understand what you might be interested in.
 - Information about the post. These are signals both about how popular a post is – think how many people have liked it and how quickly people are liking, commenting, sharing and saving a post – and more mundane information about the content itself, like when it was posted, and what location, if any, was attached to it.
 - Information about the person who posted. This helps us get a sense for how interesting the person might be to you, and includes signals like how many times people have interacted with that person in the past few weeks.
 - Your history of interacting with someone. This gives us a sense of how interested you are generally in seeing posts from a particular person. An example is whether or not you comment on each other's posts.
- **Predictions:** From there we make a set of predictions. These are educated guesses at how likely you are to interact with a post in different ways. There are roughly a dozen of these. In Feed,

the five interactions we look at most closely are how likely you are to spend a few seconds on a post, comment on it, like it, share it, and tap on the profile photo. The more likely you are to take an action, and the more heavily we weigh that action, the higher up in Feed you'll see the post. We add and remove signals and predictions over time, working to get better at surfacing what you're interested in.

We have also recently announced the publication of [22 system cards](#), providing an unprecedented level of insight into how our AI systems rank content on Facebook and Instagram. These cards provide information about how our AI systems rank content for Feed, Reels, Stories, and other surfaces; some of the predictions each system makes to determine what content might be most relevant to people; and the options available to help customize an experience on Instagram. Furthermore, in June 2023, we've also published an update explaining [Instagram Ranking](#) to help people maximize their experience and help creators understand how their content might be surfaced.

In parallel, we recently shared our [guidelines](#) for what types of content we move lower in Feed and Stories.

6. ADVERTISING TRANSPARENCY

[Explanation on how advertising works on Instagram.](#)

We have provided information about our advertising policies in Section 1 (General Observations), which sets out the types of ads content we allow or prohibit. In addition to these policies, we maintain an [Ad Library](#) that enhances advertising transparency by providing a comprehensive searchable database of all ads currently running on Instagram, along with dates the ad ran, the parameters used for targeting (e.g. age, gender, location), who was served the ad, and more.

Additionally as set out in Section 4 (Information on User Controls), when users see an ad, they can select [Why am I seeing this ad?](#) to view details on why the user is seeing a specific ad in their feed.

7. AUTOMATED MEANS FOR CONTENT MODERATION

[Any use made of automated means for the purpose of content moderation.](#)

A. Use of automated means for the purpose of own initiative and other content moderation, and purpose of those tools

As described in Section 1 (General Observations), we use technology to help us proactively detect content on our services that might be harmful and violate our Community Guidelines, so we can remove it faster. We described in Section 5 our primary content moderation automated tools.

These technologies run on accounts, posts, comments, photos, and other pieces of content uploaded to Instagram. They determine how probable or likely it is that this content goes against a certain policy, based on those signals or patterns, and if the content should be automatically removed.

B. Indicators of accuracy, error rates, safeguards

Our technology learns and improves from each human decision. Over time – after learning from thousands of human decisions – the technology gets better. When reviewing violating content, review teams manually label the policy guiding their decision, which means that they mark or “label” the relevant policy that the content, account, or behavior violates. This labeling of data helps us improve the quality of our algorithms that proactively detect and remove harmful content, accounts, and behavior.

To ensure and improve the quality, i.e., how accurate the technologies mentioned above are in enforcing Community Guidelines and other policy violations, there are ongoing quality evaluation processes in place. Meta uses overlapping techniques and systems for maintaining a high overall accuracy for our automated content moderation.

Prior to fully launching any new rate limit (speed limit), matching technology or machine learning (ML), we use the technology to only log how the technology would have behaved instead of immediately acting. We then use human reviewers to assess the accuracy against current content, behavior, or accounts, rather than just historical ones, as we did during the technology’s training. After launching rate limits, matching technologies, or machine learning, we monitor the volumes of actions and appeals by the user who posted the content as well as the rate at which appeals are granted. If any of the metrics we monitor are abnormal, our engineering teams may investigate.

For each primary form of automation technology, the investigation of abnormal metrics can vary. With rate limits, engineers typically reevaluate if the limit is preventing bot behavior. For our matching technologies, if an entry in our list of previously identified instances of policy violations has abnormal signals, we will re-review the entry to confirm it continues to go against our policies. Similarly, if one of our machine learning classifiers has abnormal signals, we will either send a sample

of the machine learning classifier's recent results to human labeling to confirm the accuracy rate or deprecate the classifier if abnormal signals indicate a clear breakage.

In addition, many of our machine learning classifiers are automatically reassessed for accuracy after each human review. This classifier reassessment is an example of the general feedback loop between human review and technology. The content labeling decisions taken by human reviewers are used to train and refine our technology. As a part of this process, the review teams manually label the policy guiding their decision, i.e., they mark the policy that the content, account, or behavior violates. This helps to improve the quality of our machine learning algorithms and our lists of known policy-violating content used by our matching technology. To maintain quality control in all of these decisions, we regularly audit random samples of decisions taken by the algorithm and our content reviewers and measure them against our expectations for policy enforcement. In the context of automation relating to language, some automation is developed to support specific languages whilst others are language agnostic.

8. PERSONS IN CHARGE OF CONTENT MODERATION

The human resources that Meta dedicates to content moderation in respect of the service offered in Türkiye, their qualifications and linguistic expertise and training and support given to such staff.

As described in Section 1 (General Observations), Meta uses human resources dedicated to content moderation, including for compliance with the obligations set out in Articles 8 and 8/A and Articles 9 and 9/A of the Law No. 5651.

A. Qualifications And Training

Human reviewers come from different backgrounds, reflect our diverse community, and include experts in enforcement in policy areas such as child safety, hate speech, and counterterrorism. Human reviewers undergo extensive training when they join and are regularly trained and tested beyond this initial training, with specific examples, such as on how to uphold the Community Guidelines and take the correct action on a report. We also do our own proactive audits, where we conduct re-reviews that help us figure out if we are getting it right.

Human reviewers who review content alleged to be illegal receive distinct training based on the nature of their respective work. Every member receives several weeks of training focused heavily on operational proficiency and in preparation for processing such content. For example, the reviewers

who review content for defamation receive training specifically on assessing defamation, including background material on the law.

B. Support

We recognise that reviewing content can be challenging work. Keeping people safe online sometimes means review teams have to look at content that may be objectionable or graphic. We respect the difficulty of this work and work with industry leading vendors to ensure reviewers have access to the resources they need to do their job and support their health. There is a robust and diverse program to support human reviewers. Our vendor contracts mandate high quality support in a variety of areas, including pay, benefits, work environment, and wellbeing and psychological support. The assistance model depends on what type of content reviewers work on. Such assistance, for example, can take the form of psychological support, including individual and group sessions, and a 24/7 independent support program which includes a range of offerings that include clinical services.

C. Volume of human resources dedicated to content moderation and linguistic capabilities

The team working on safety and security is made up of around 40,000 people. About 15,000 of those are content reviewers, and they include a mixture of full-time employees, contractors, and outsourced support. We partner with companies to help with content review, which allows us to scale globally with coverage across time zones, languages, and markets.

This content review team is global and reviews content 24/7 in over 80 languages, including Turkish. The team includes reviewers with language expertise to enforce our policies in cases where certain words or content require additional contextual understanding. Not all content requires language expertise; for example, nudity and sexual activity is language agnostic. For this type of content, Meta has a global pool of content moderators who review these types of reports. Currently, there are around 2,500 language agnostic reviewers globally. These reviewers also review Turkish content that falls under this language category.